

Module 01: Creating Azure App Service Web Apps



Topics

- Azure App Service core concepts
- Creating an Azure App Service Web App
- Configuring and Monitoring App Service apps
- Scaling App Service apps
- Azure App Service staging environments

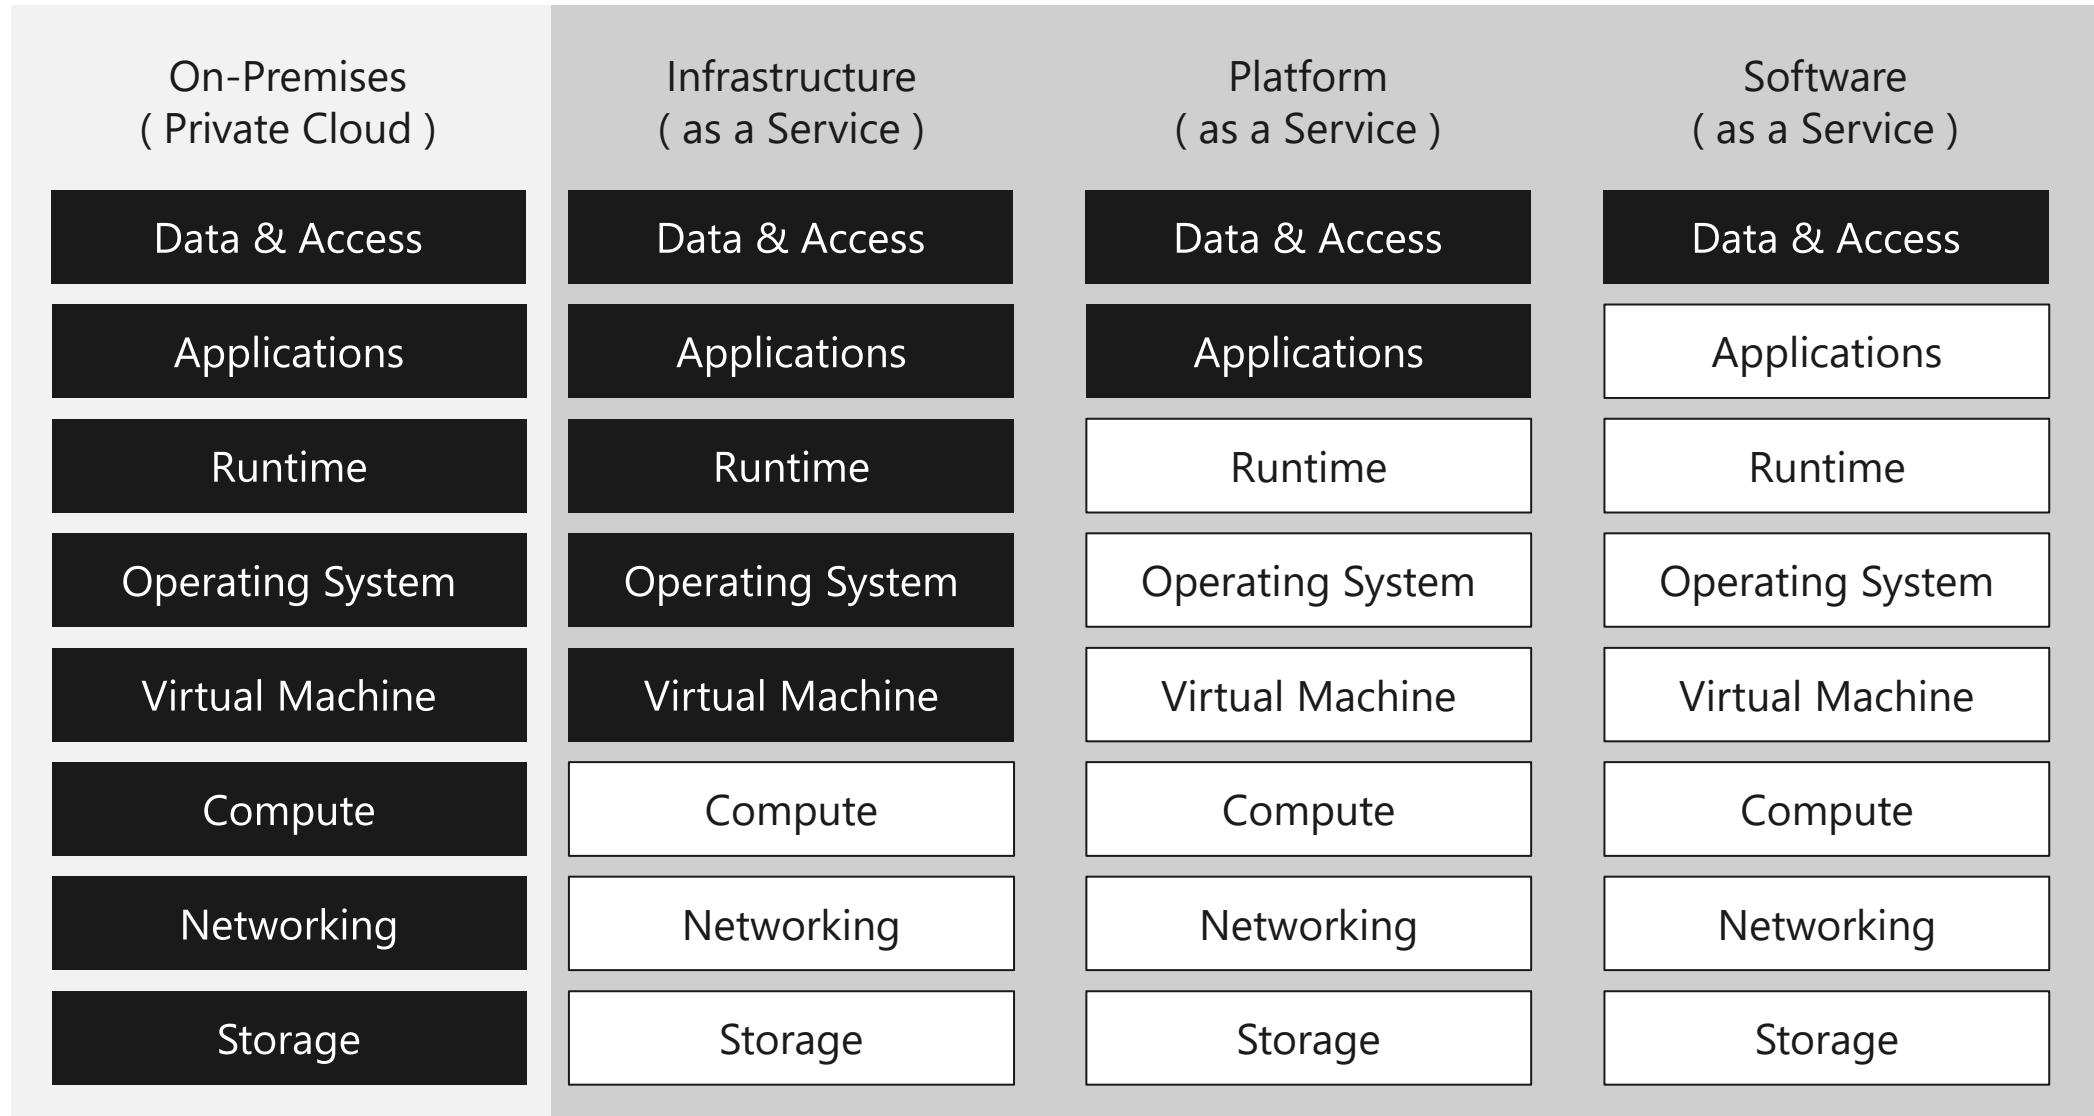
Lesson 01: Azure App Service core concepts



You Manage

Azure shared responsibility model

Cloud Provider Manages



App Service

- Service for hosting web applications, REST APIs, and mobile backends can be developed in many of the following languages:

.NET

Java

Ruby

Node.JS

PHP

Python

- Applications can execute and scale in a fully managed, sandbox environment

Web Apps

- Scalable hosting for web applications:
 - Provides a quick way to host your web application in the cloud
 - Allows you to scale your web app without being required to redesign for scalability
 - Integrates with Visual Studio
 - Provides an open platform for many different programming languages
- Advantages:
 - Near instant deployment
 - SSL and Custom Domain Names available in some tiers
 - WebJobs provide background processing for independent scaling
 - Can scale to larger machines without redeploying applications

Key features of App Service Web Apps

- Multiple languages and frameworks:
 - First-class support for Microsoft ASP.NET, Java, Ruby, Node.js, PHP, or Python
- DevOps optimization:
 - Continuous integration and deployment with Visual Studio Team Services, GitHub, Bitbucket, Docker Hub, or Azure Container Registry
- Global scale with high availability:
 - Scale up or out manually or automatically. Host anywhere in the Microsoft global datacenter infrastructure
- Connections to SaaS platforms and on-premises data:
 - More than 50 connectors for enterprise systems (such as SAP), SaaS services (such as Salesforce), and internet services (such as Facebook)

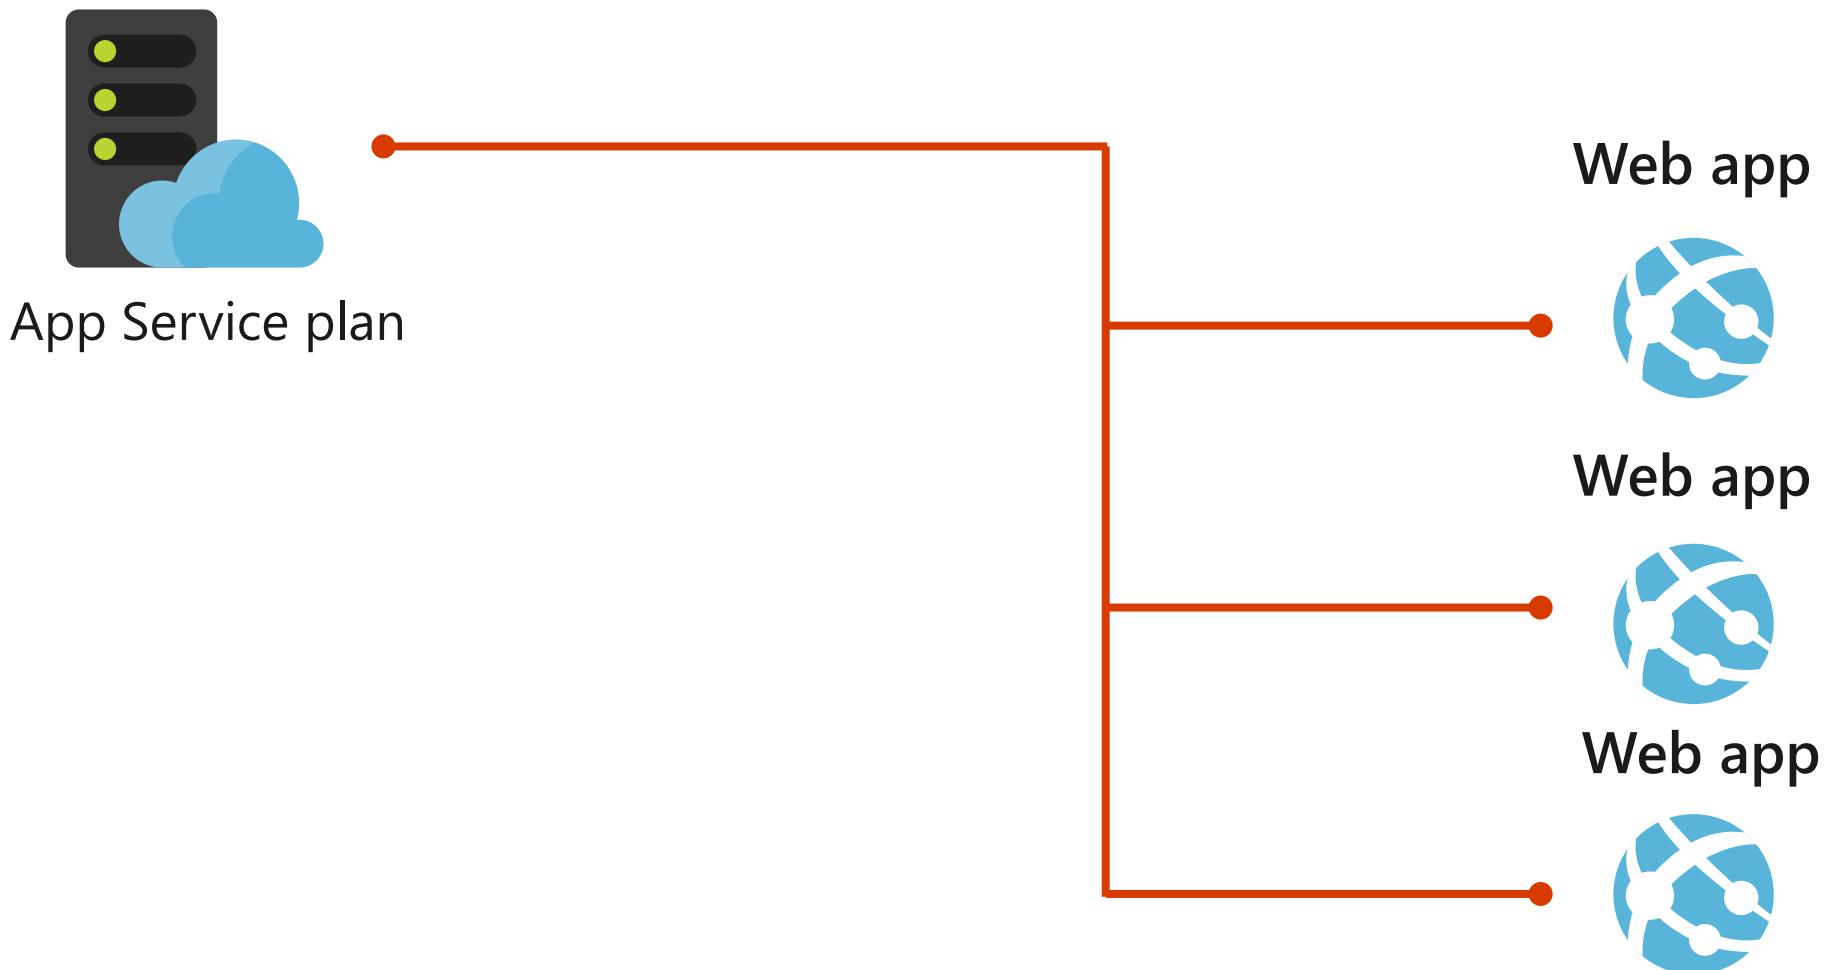
Key features of App Service Web Apps (cont.)

- Security and compliance:
 - App Service is ISO, SOC, and PCI compliant
- Application templates:
 - Templates in the Azure Marketplace, such as WordPress, Joomla, and Drupal
- Visual Studio integration:
 - Streamline the work of creating, deploying, and debugging
- API and mobile features:
 - Turn-key Cross-Origin Resource Sharing (CORS) support for RESTful API scenarios, and enables authentication, offline data sync, push notifications, and more
- Serverless code:
 - Run code on-demand without having to explicitly provision or manage infrastructure

App Service plans

- App Service plans can logically group apps within a subscription:
 - Characteristics such as features, capacity, and tiers are shared among the website instance in the group
 - The App Service plan is the unit of billing in most cases
- Multiple App Service plans can exist in a single Resource Group and multiple apps can exist in a single App Service plan

App Service plans (continued)



Lesson 01: Creating an Azure App Service Web App

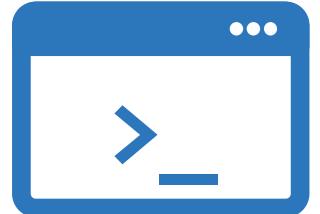


Demonstration: Creating a web app by using the Azure portal



Creating a web app with Azure command-line interface (CLI) (continued)

```
# generate a unique name and store as a shell variable  
webappname=mywebapp$RANDOM  
  
# create a resource group  
az group create --location westeurope --name myResourceGroup  
  
# create an App Service plan  
az appservice plan create --name $webappname --resource-group myResourceGroup --sku  
FREE  
  
# create a Web App  
az webapp create --name $webappname  
  --resource-group myResourceGroup  
  --plan $webappname
```

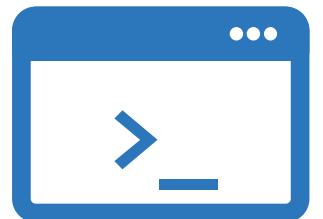


Deploying a web app with Azure CLI

```
# store a repository url as a shell variable
gitrepo=https://github.com/Azure-Samples/php-docs-hello-world

# deploy code from a Git repository
az webapp deployment source config --name $webappname --resource-group myResourceGroup
--repo-url $gitrepo --branch master --manual-integration

# print out the FQDN for the Web App
echo http://$webappname.azurewebsites.net
```



Creating a Web App with Azure PowerShell

Command	Notes
New-AzResourceGroup	Creates a resource group in which all resources are stored
New-AzAppServicePlan	Creates an App Service plan
New-AzWebApp	Creates an Azure Web App
Set-AzResource	Modifies a resource in a resource group

Creating an App Service plan with Azure PowerShell

```
# Create variables for the repository URL, Web App name and location
$gitrepo="https://github.com/Azure-Samples/app-service-web-dotnet-get-started.git"
$webappname="mywebapp$(Get-Random)"
$location="West Europe"

# Create new resource group
New-AzResourceGroup -Name myResourceGroup -Location $location

# Create new App Service plan
New-AzAppServicePlan -Name $webappname -Location $location -ResourceGroupName
myResourceGroup -Tier Free
```



Creating a Web App with Azure PowerShell

```
# Create new Web App
New-AzWebApp -Name $webappname -Location $location -AppServicePlan $webappname -
ResourceGroupName myResourceGroup

# Create deployment resource manually using ARM
$PropertiesObject = @{
    repoUrl = "$gitrepo";
    branch = "master";
    isManualIntegration = "true";
}
Set-AzResource -PropertyObject $PropertiesObject -ResourceGroupName myResourceGroup -
 ResourceType Microsoft.Web/sites/sourcecontrols -ResourceName $webappname/web -
 ApiVersion 2015-08-01 -Force
```



Demonstration: Creating a static HTML web app by using Azure Cloud Shell

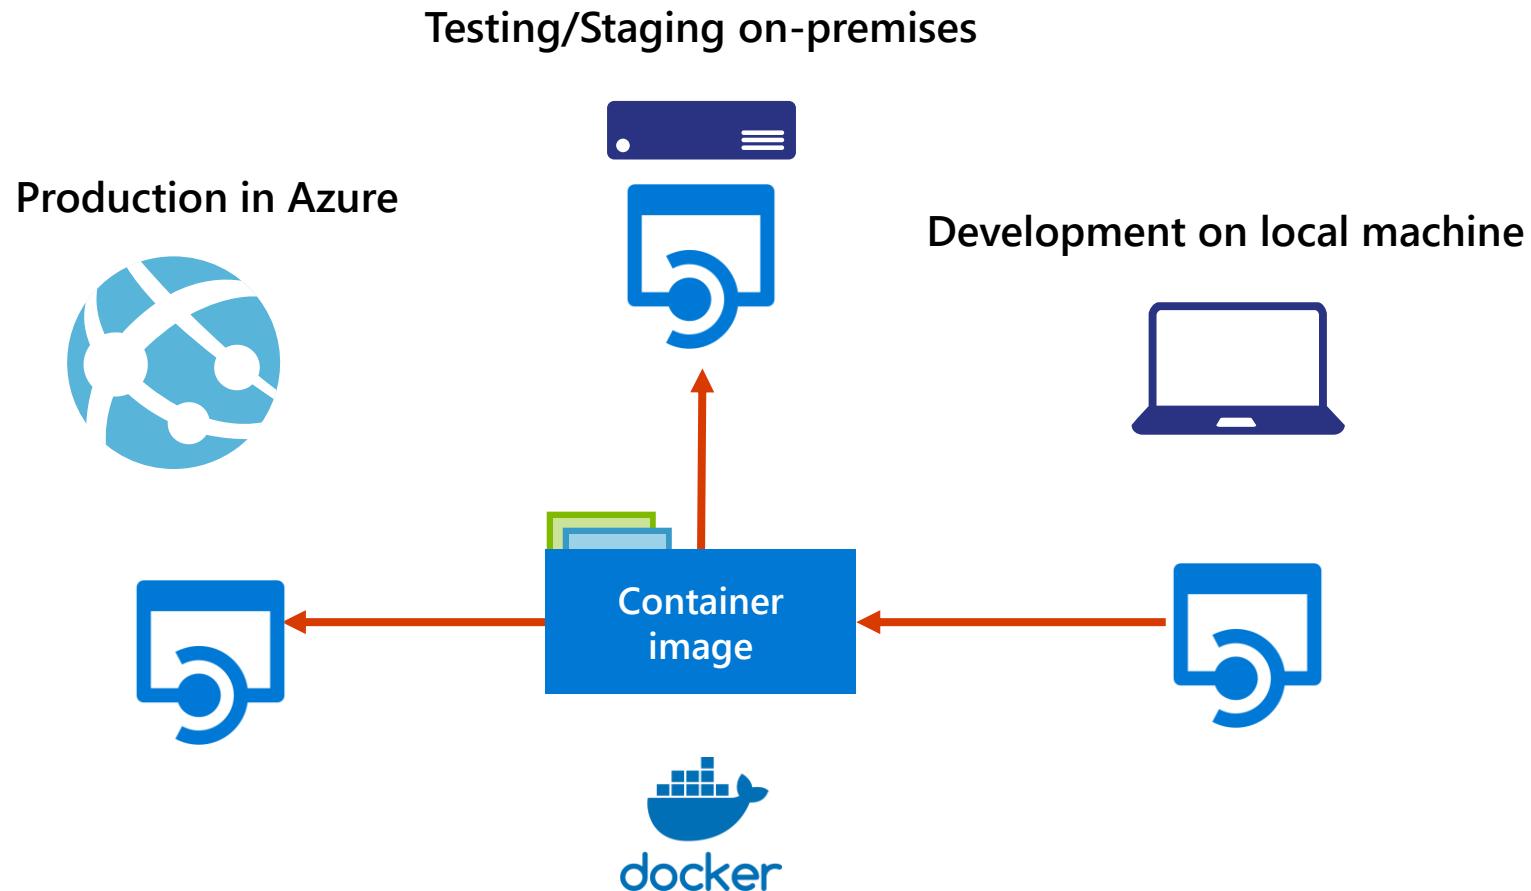


App Service on Linux

Why Linux?

- Many application stacks are optimized for Linux:
 - Ruby/Rails, PHP, Node, and others
 - Often, better tools are available on Linux for these stacks
- New and upcoming frameworks are built for Linux first and then Windows
- Portability of Docker containers
- Linux is at the forefront of innovations in nano and microservice architecture

Docker in App Service on Linux



Web apps for Linux containers

Deploy applications and solutions that are containerized directly to App Service Web Apps:

- Simplifies deployment
- Matches the already popular container workflow using:
 - CI/CD with Docker Hub, Azure Container Registry, or GitHub
- Compatible with existing App Service features:
 - Auto-scale, Deployment Slots, and others

Web apps for Linux containers (continued)

Containers can be sourced from your existing registries:

- Docker Hub:
 - Deploy images already shared on Docker Hub
 - Deploy the most popular official images
 - Private images are available on Docker Hub
- Azure Container Registry:
 - Managed service for hosting Docker images
 - Can deploy to Docker Swarm, Kubernetes, or Web App for Containers

Demonstration: Creating Web Apps with a local Git deployment source



Lesson 03: Configuring and Monitoring App Service apps



App Service settings

- Overrides settings in Web.config or appsettings.json
- Hidden by default in Azure portal
- You can configure:

Application
settings

Connection
strings

Default
documents

Path mappings

Language
stack (app
runtime)

Custom
containers

Default documents

- Only available for Windows apps
- List of documents to show when navigating to a directory on the web server:
 - First matching file is used
- Alternative to building a custom module

Path mappings

- Windows:
 - Custom IIS handler mappings
 - Virtual applications/directories
- Containerized:
 - Custom-mounted storage

Updating app runtimes

```
az webapp config set --linux-fx-version "DOTNETCORE|3.1" --resource-group <groupname> --name <appname>
```

Update .NET version

```
az webapp config set --php-version 7.0 --resource-group <groupname> --name <appname>
```

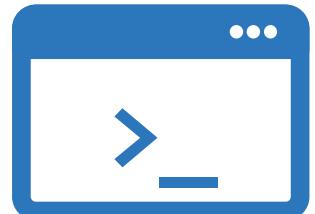
Update PHP version

```
az webapp config set --python-version 3.4 --resource-group <groupname> --name <appname>
```

Update Python version

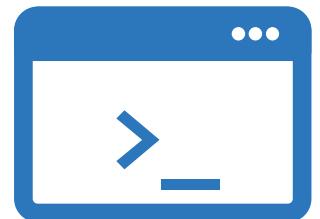
```
az webapp config set --resource-group <groupname> --name <appname> --java-version 1.8 --java-container Tomcat --java-container-version 9.0
```

Update Java version



Updating app runtimes (Node.js)

```
az webapp config appsettings set `  
  --resource-group <groupname> `  
  --name <appname> `  
  --settings WEBSITE_NODE_DEFAULT_VERSION=8.9.3
```



Editing bulk settings

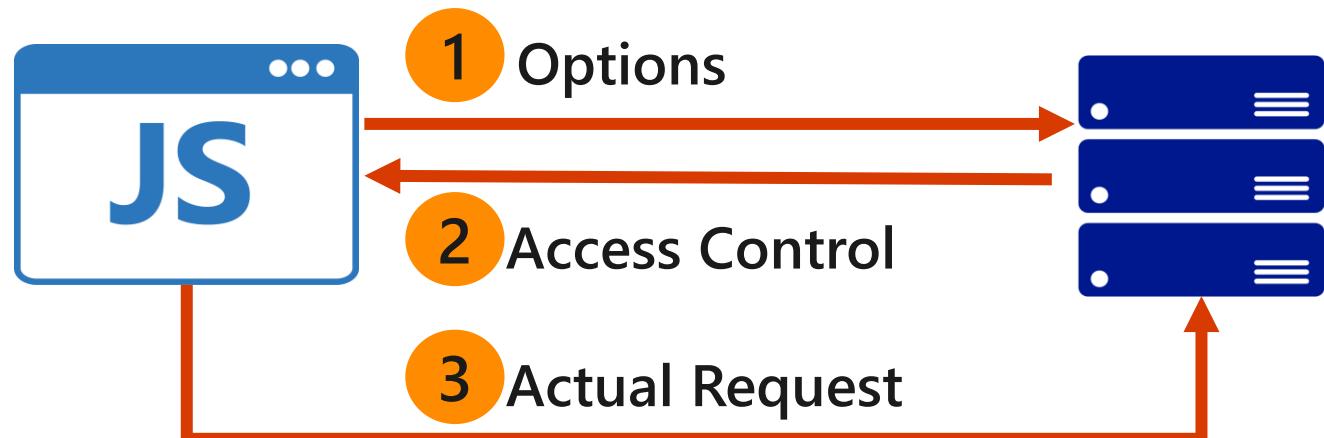
- Allows editing App Settings in Bulk

The screenshot shows the Azure portal's Application settings page for a web app. On the left, the portal interface displays application settings like ANCM_ADDITIONAL_ERROR_PAGE_LINK, APPINSIGHTS_INSTRUMENTATIONKEY, and others. On the right, a large red arrow points from the portal UI to a block of JSON code, highlighting the "slotSetting: false" property which indicates that these settings apply to all slots.

```
[{"name": "ANCM_ADDITIONAL_ERROR_PAGE_LINK", "value": "https://bicowartung.scm.azurewebsites.net/detectors", "slotSetting": false}, {"name": "APPINSIGHTS_INSTRUMENTATIONKEY", "value": "3aa11fae-ef13-4f11-9941-33ada55d9aa8", "slotSetting": false}, {"name": "APPINSIGHTS_PROFILERFEATURE_VERSION", "value": "disabled", "slotSetting": false}, {"name": "APPINSIGHTS_SNAPSHOTFEATURE_VERSION", "value": "disabled", "slotSetting": false}, {"name": "ApplicationInsightsAgent_EXTENSION_VERSION", "value": "~2", "slotSetting": false}, {"name": "DiagnosticServices_EXTENSION_VERSION", "value": "1.0.0", "slotSetting": false}]
```

CORS

- Mechanism for servers to indicate that they support cross-site requests
 - Servers can specify:
 - Allowed HTTP verbs
 - Allowed origins
 - Allowed headers
 - Directly supported by API Apps



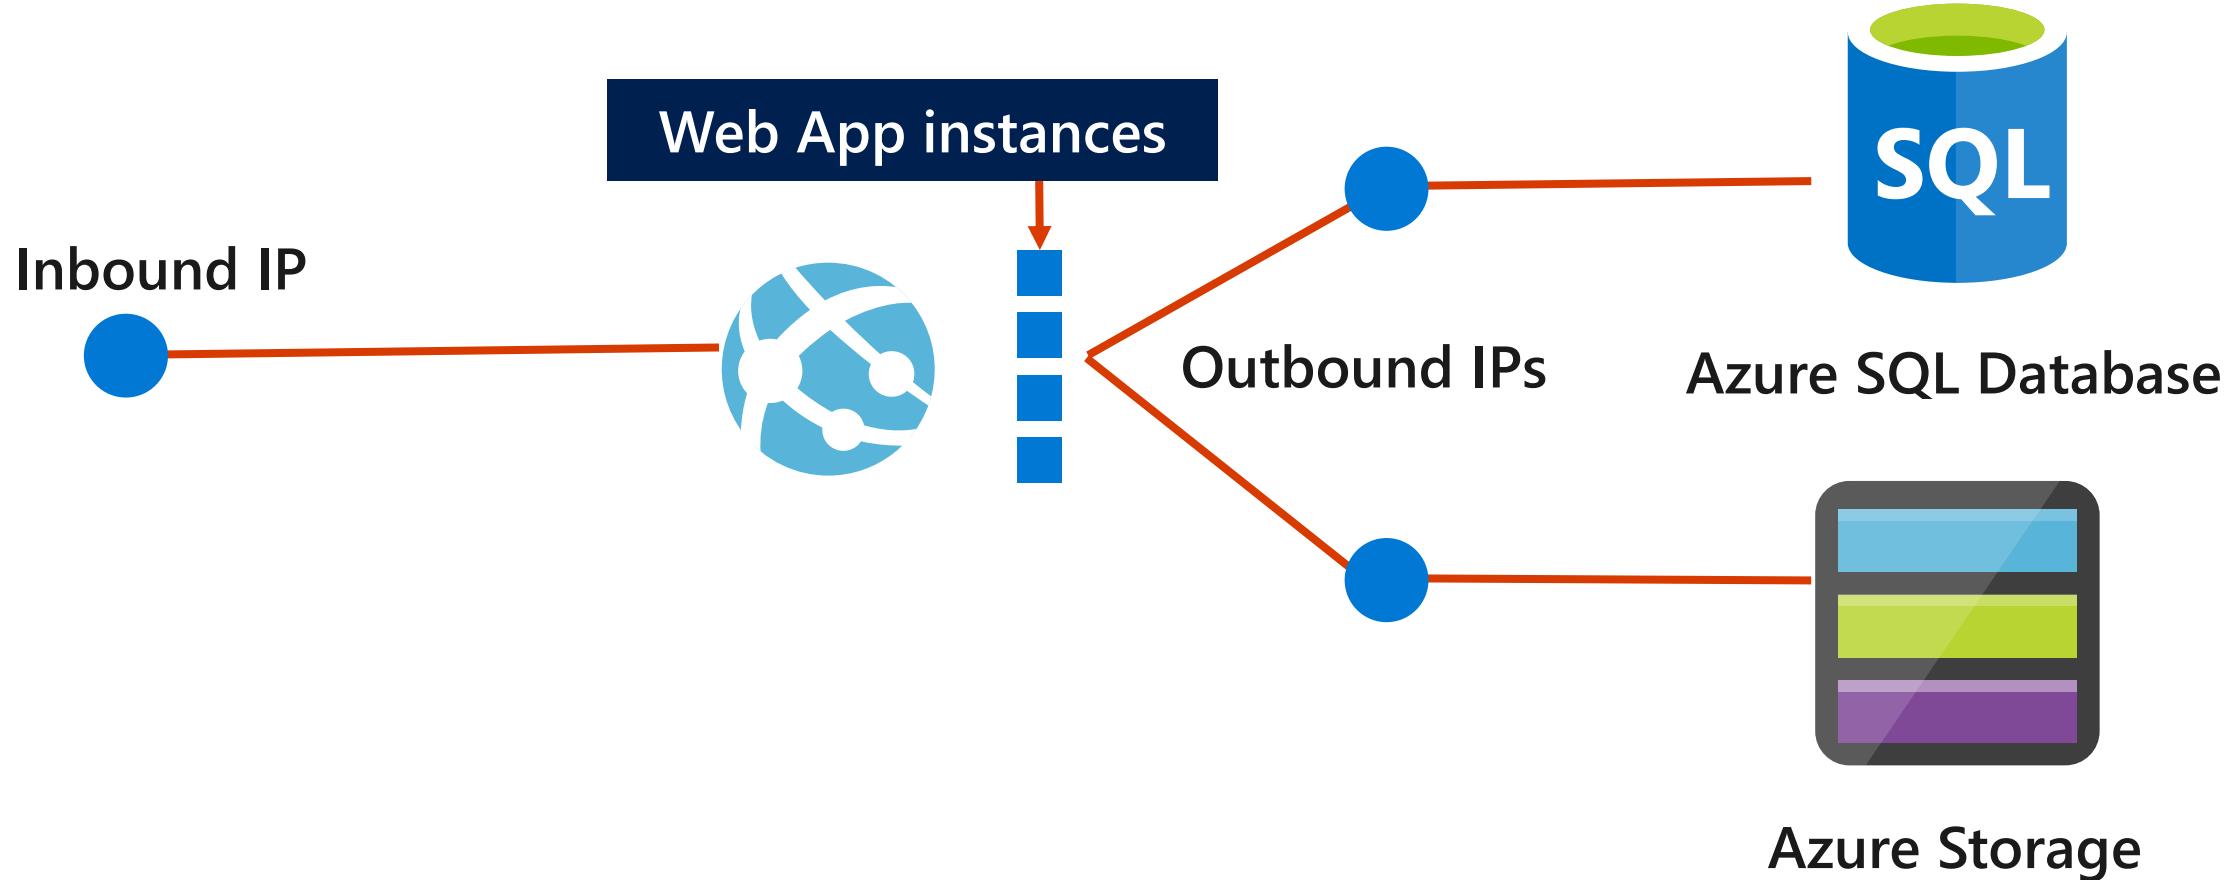
OS and runtime patching

- OS and application stack are managed by Azure on your behalf
- Monthly OS patching:
 - Physical servers
 - Guest virtual machines
- Stable versions of application runtimes are periodically added to App Services:
 - Some are installed side by side, while others replace existing versions
 - You can manually migrate from one application runtime to another

Inbound and outbound IP addresses

- Each app has a single inbound IP address:
 - Regardless of scale-out quantity
- Inbound IP address can change:
 - Delete an app and re-create it in a new resource group
 - Delete the last app in a resource group + region combination and re-create it
 - Delete an existing SSL binding
- You can opt to use a static inbound IP
- Each app has a set number of outbound IP addresses:
 - The set and quantity changes as you scale your app between tiers

Outbound IP addresses

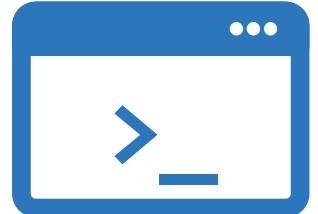


When IP addresses change

- Inbound:
 - When you delete an app and recreate it
 - When you delete the last app in a resource group
 - When you delete an existing SSL binding
- Outbound:
 - Why you scale from a lower tier (Basic, Standard, Premium) to the Premium V2 tier

Find outbound IP addresses

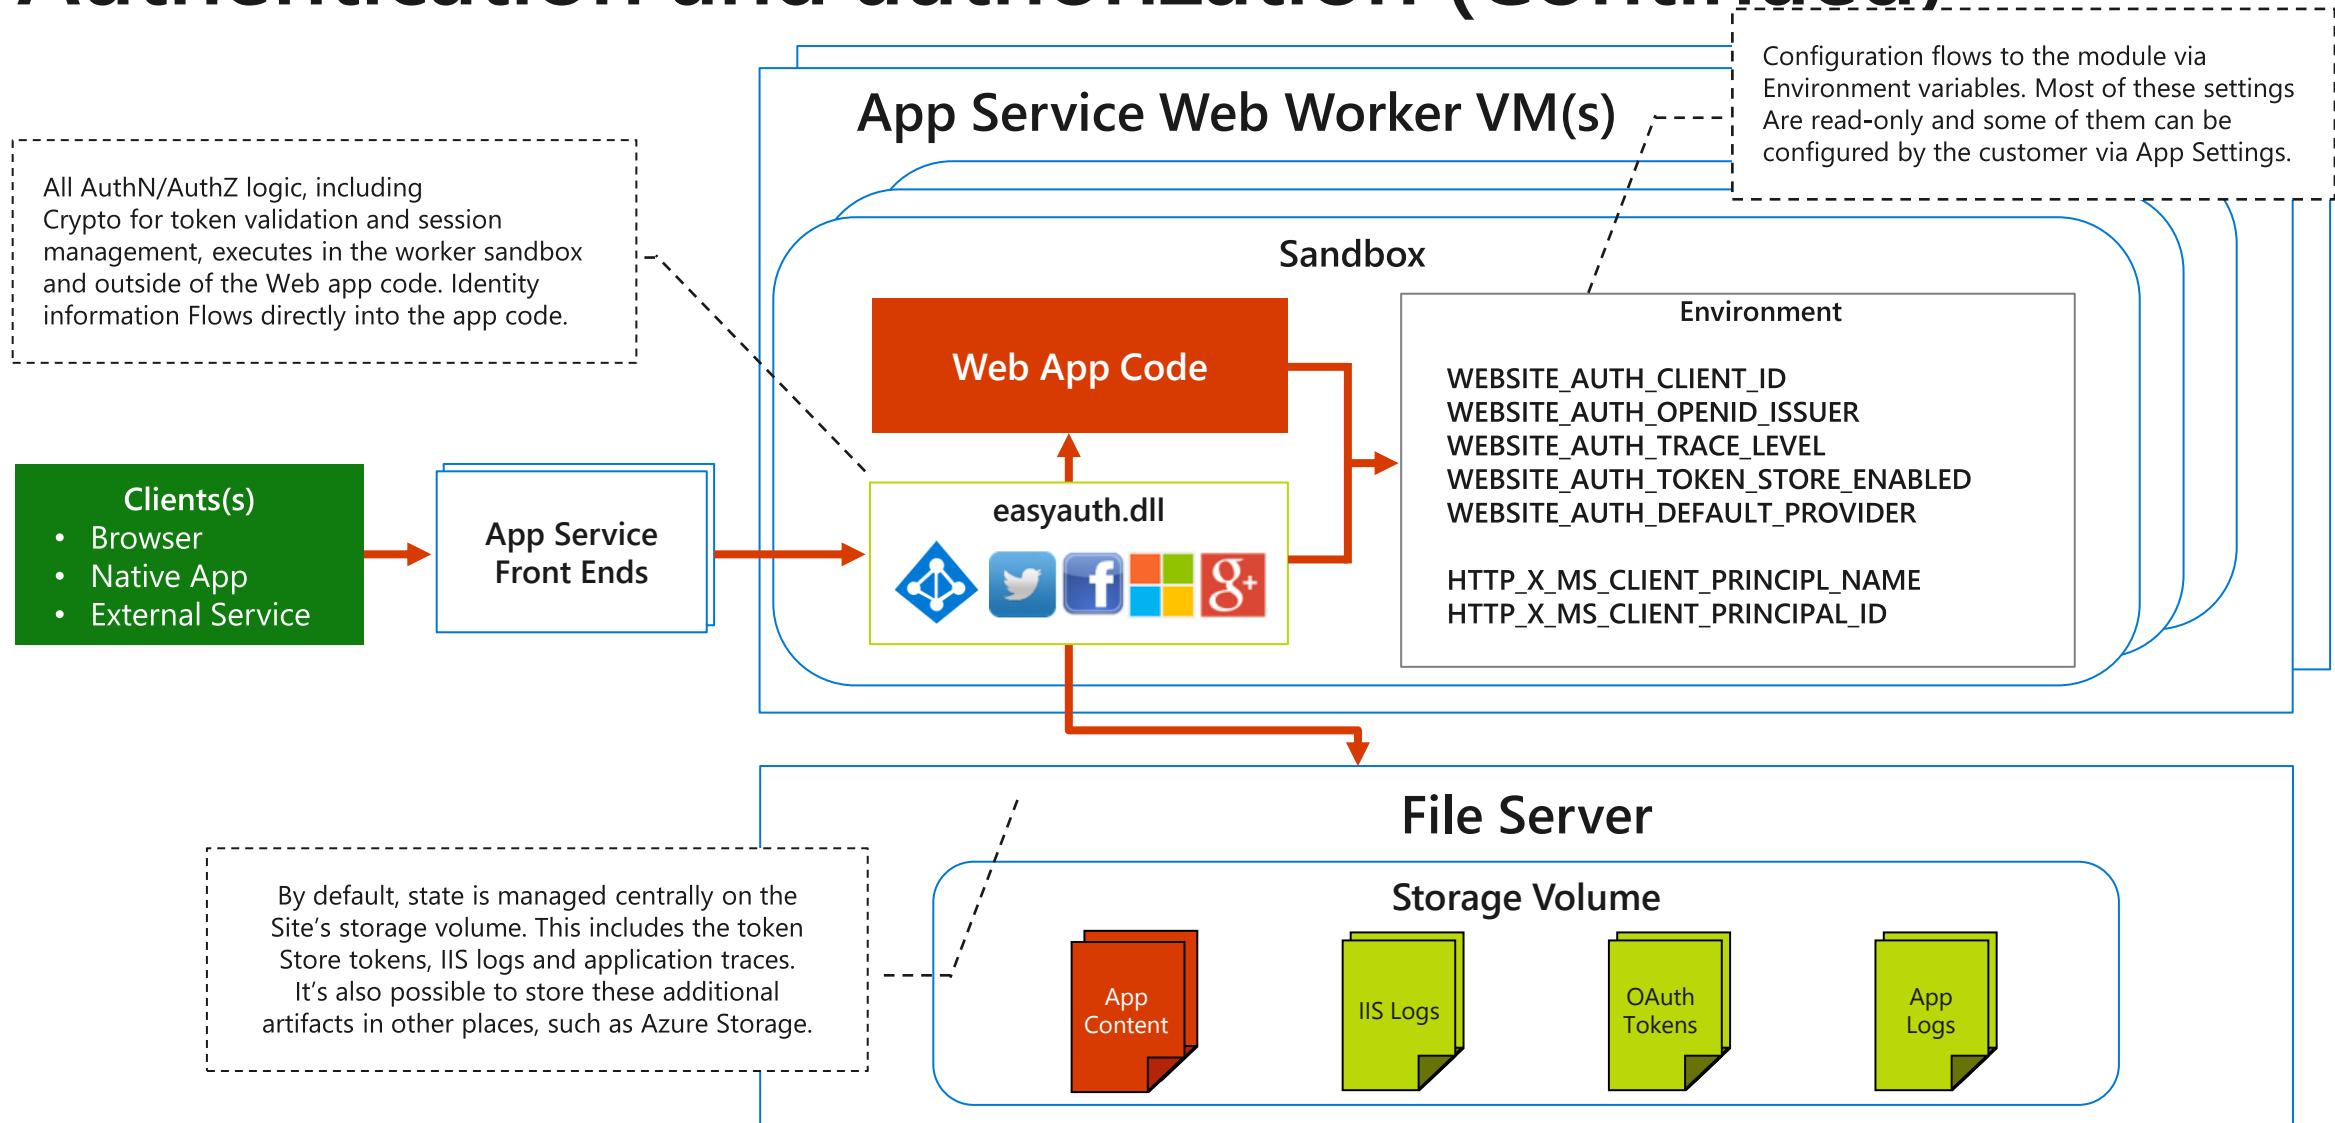
```
# find Outbound IP address  
az webapp show --resource-group <group_name> --name <app_name> --query  
outboundIpAddresses --output tsv  
  
# find all possible IP addresses (regardless of tier)  
az webapp show --resource-group <group_name> --name <app_name> --query  
possibleOutboundIpAddresses --output tsv
```



Authentication and authorization

- Built-in authentication and authorization support:
 - No extra code required to make use of these features
- User claims are made available to code:
 - If you wish to enhance the authentication support, you can use your existing code with popular identity frameworks:
 - ASP.NET Identity
 - PHP server variables
- Built-in token store
- Logging and tracing enabled for authentication events
- Support for popular identity providers:
 - Azure Active Directory (Azure AD), Microsoft accounts, Facebook, Google, Twitter, more...

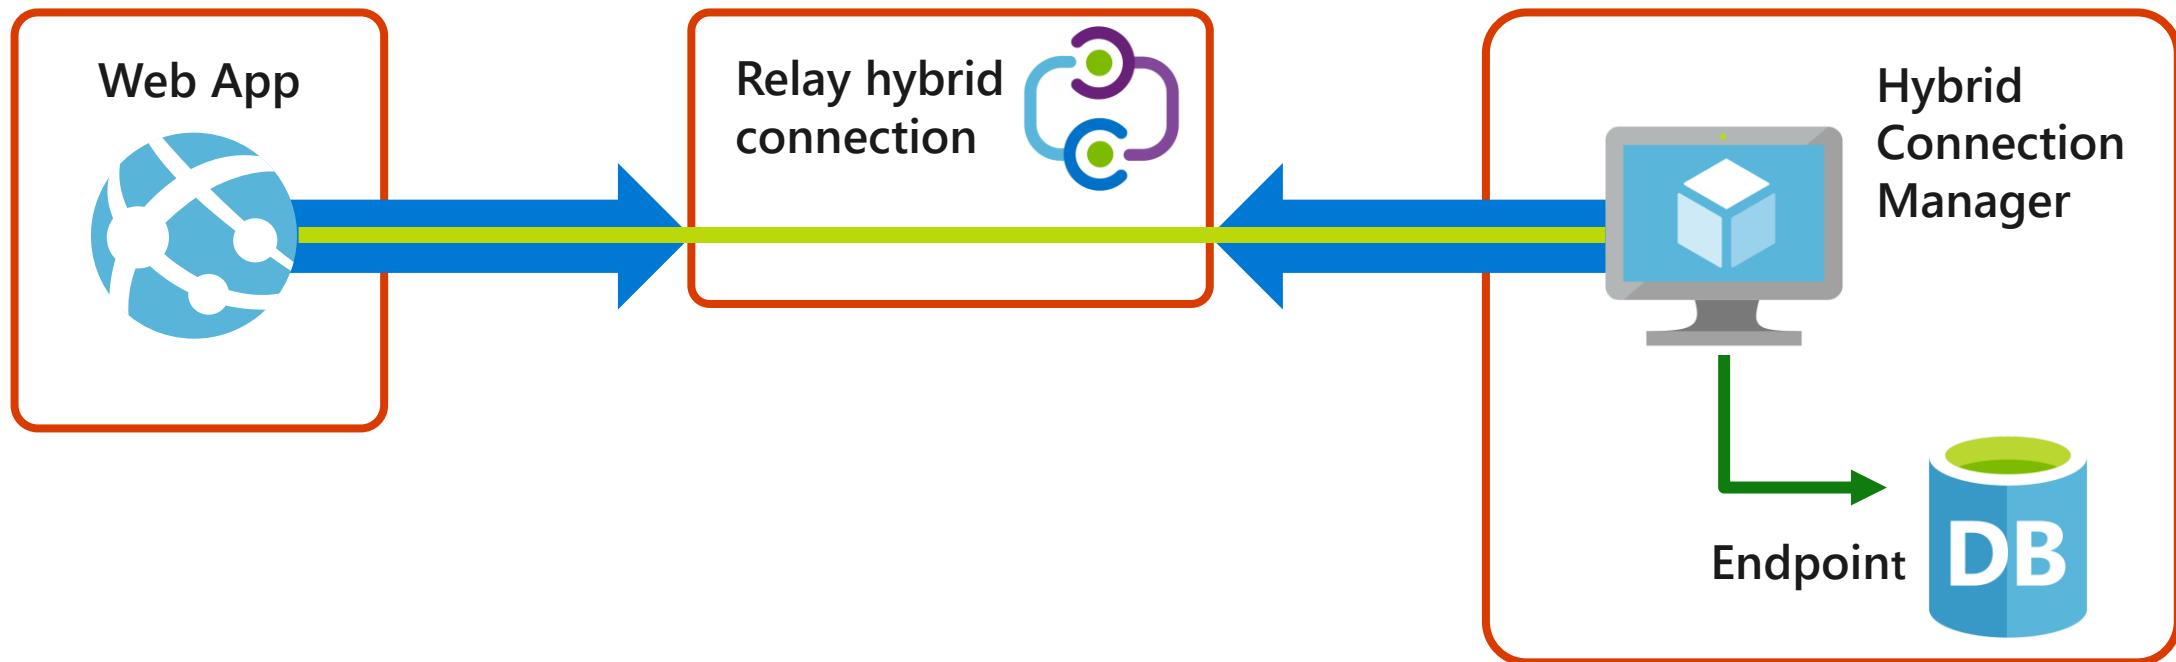
Authentication and authorization (Continued)



Azure App Service Hybrid Connections

- Enables access to resources in other networks:
 - Any operating system and any application
 - Hosted in other cloud networks, local networks, or even a specific machine
- Correlates to a single TCP host and port combination
- Service that is available in more than just App Service
- Benefits:
 - Doesn't require internet-facing endpoint
 - Web Apps can access on-premises systems securely
 - Does not require firewall changes in most scenarios
 - Framework and operating system agnostic

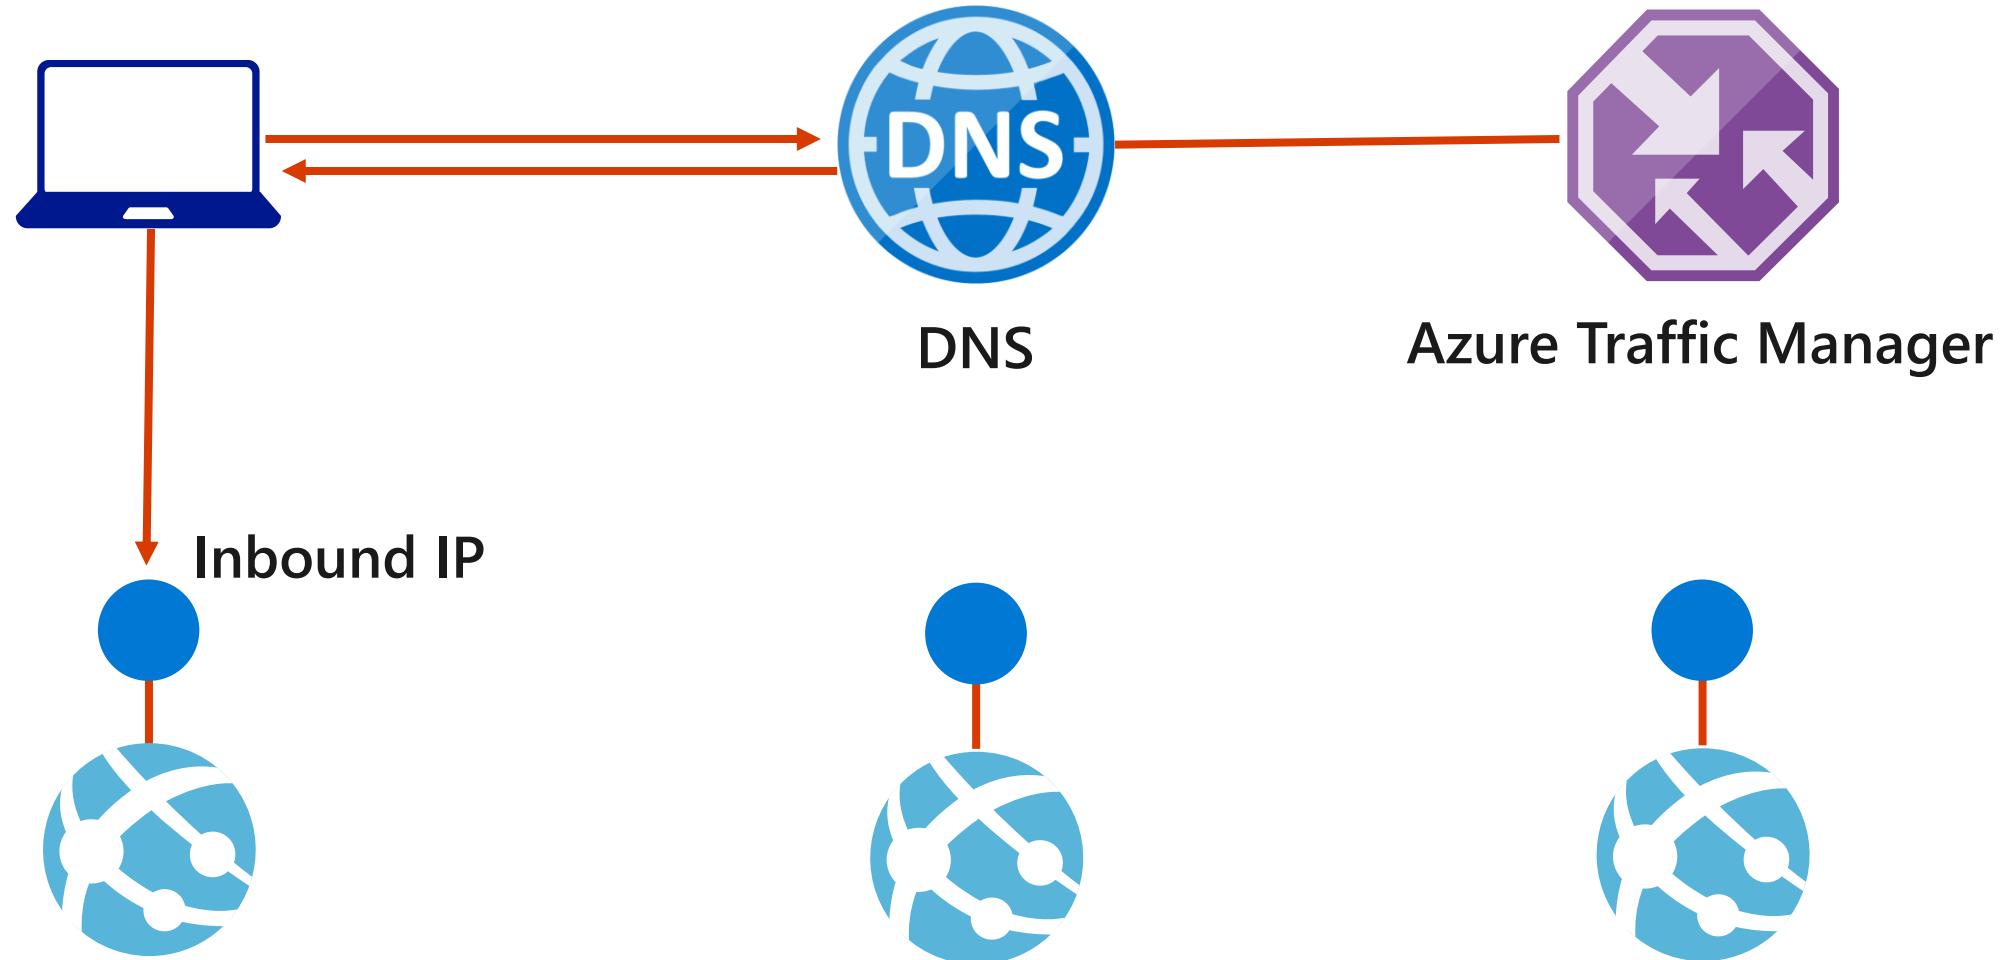
Azure App Service Hybrid Connections example



Controlling traffic by using Azure Traffic Manager

- Routes requests from clients to apps in Azure
- Keeps track of app status (running, stopped, deleted):
 - Will automatically route traffic away from an unavailable app
- Configured by using profiles:
 - Stores the routing method for requests
 - Stores a list of endpoints (apps) to route requests to
 - Stores information about endpoint status

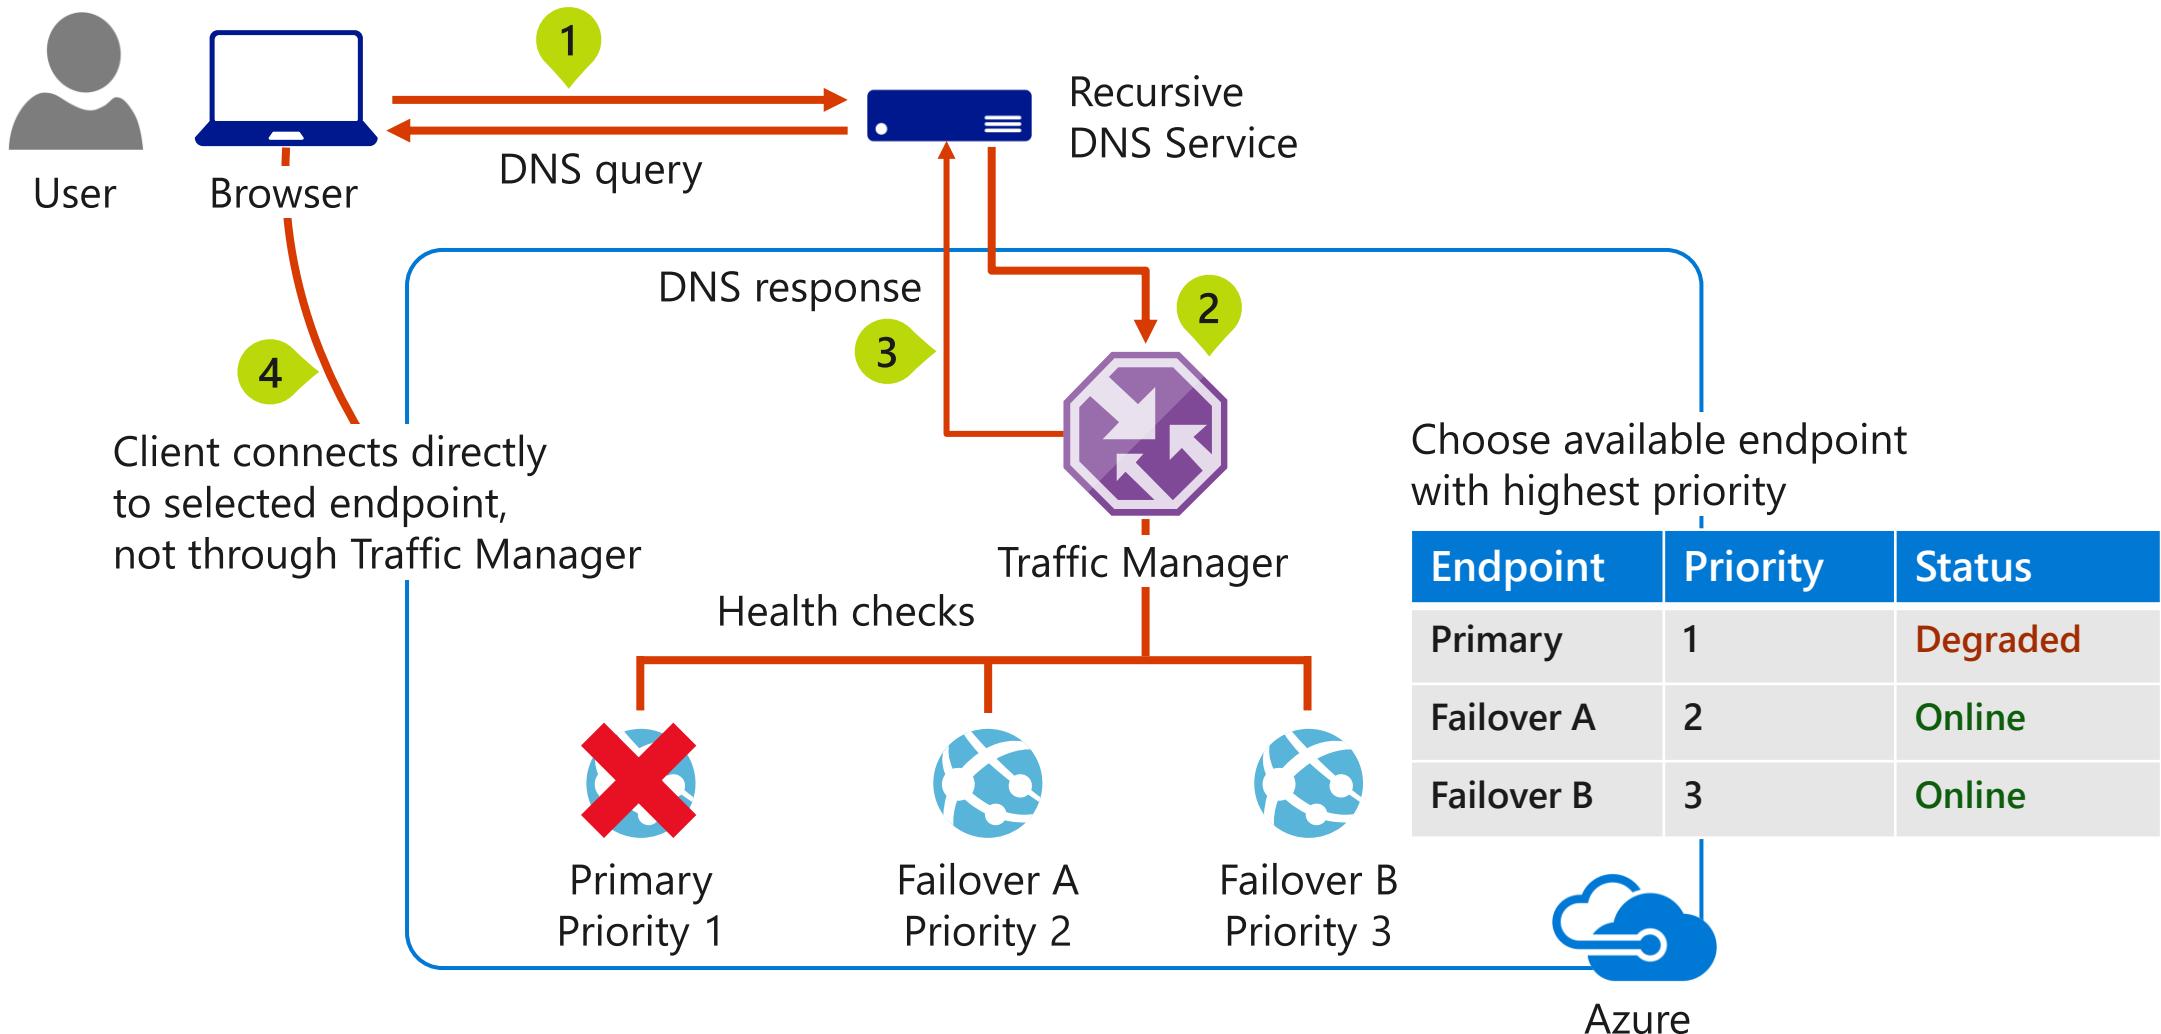
Azure Traffic Manager and Web Apps



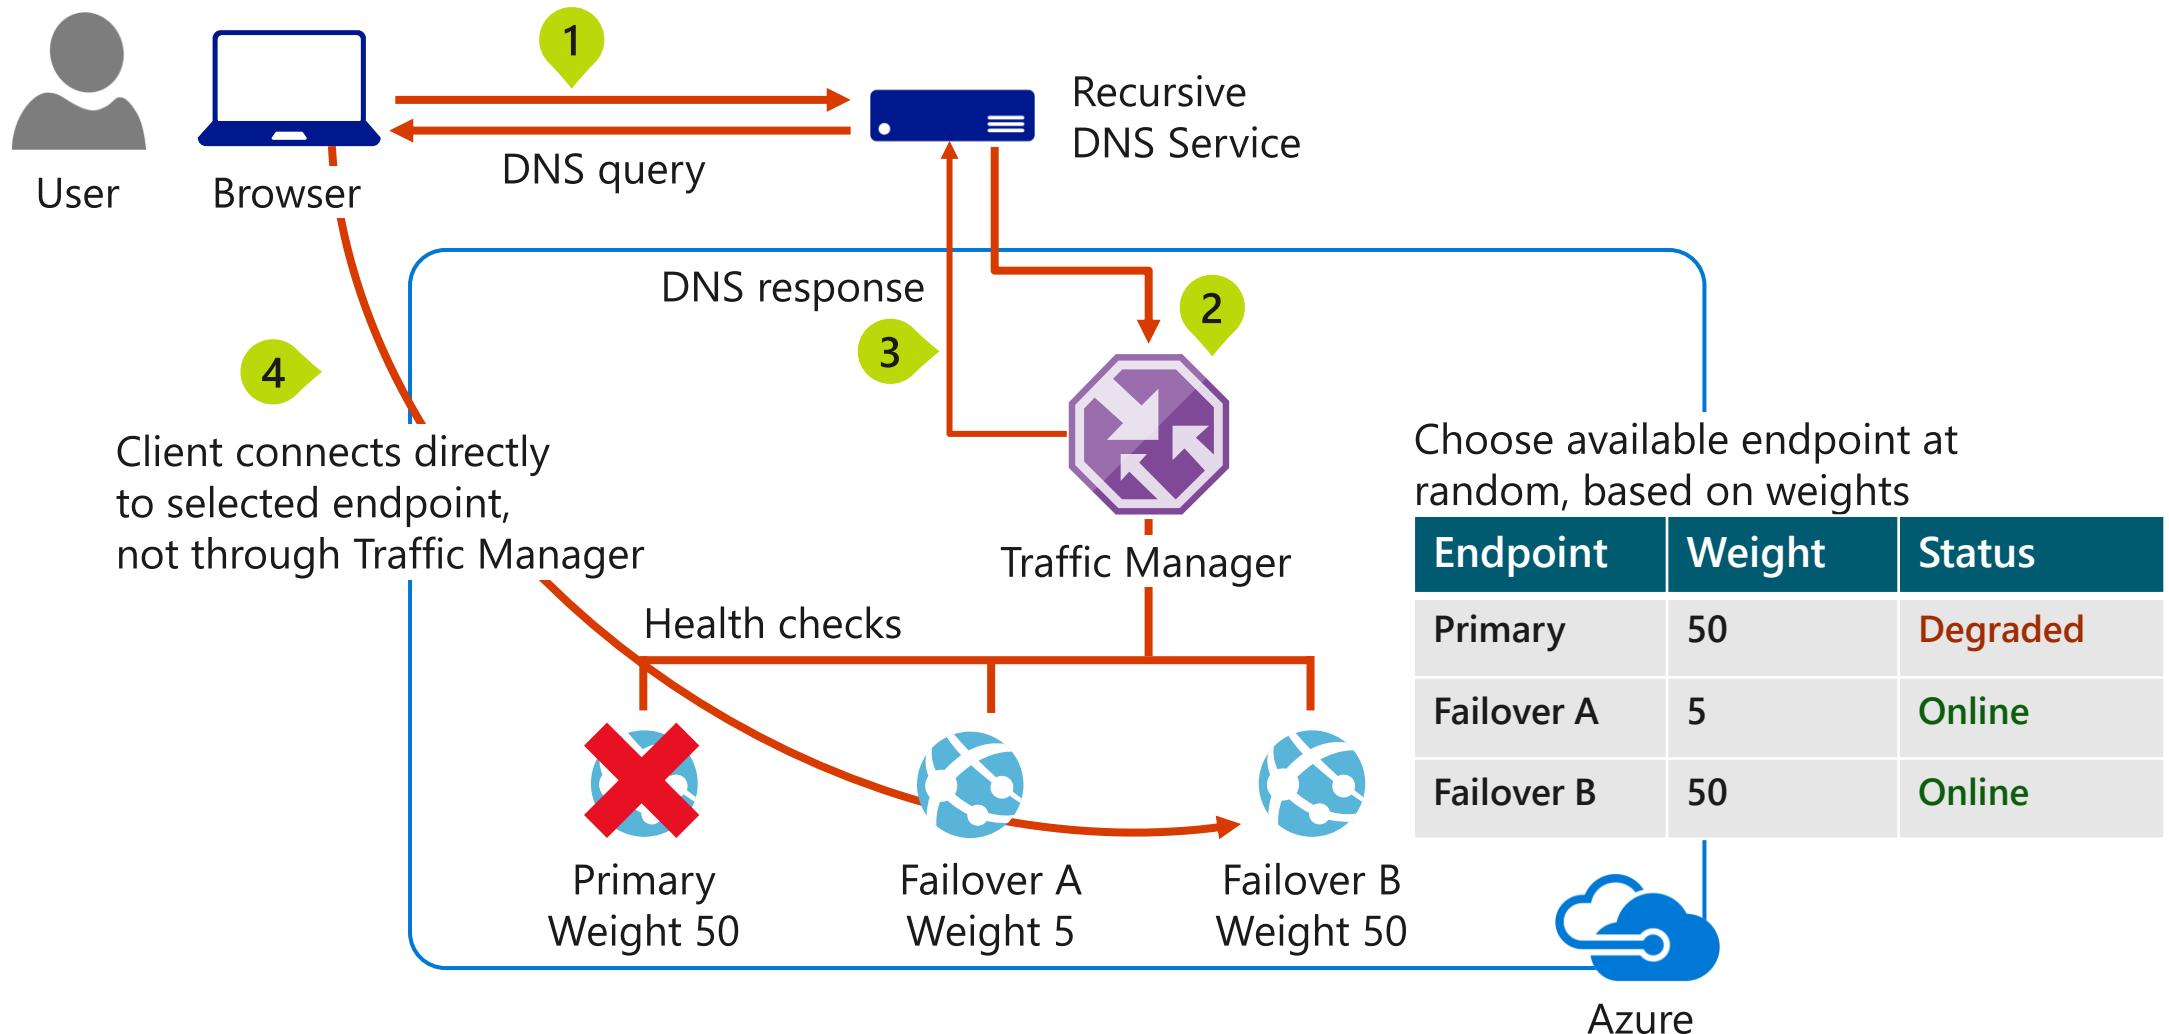
Azure Traffic Manager routing methods

- Priority:
 - Distribute users to a specific app
 - In case of failure, route users to backup apps based on a priority scheme
- Weighted:
 - Distribute traffic across apps according to weights that you define
 - Your weight definition could potentially distribute users evenly
- Performance:
 - Route users to the “closest” app location based on latency
- Geographic:
 - Route users to specific app locations based on their current location

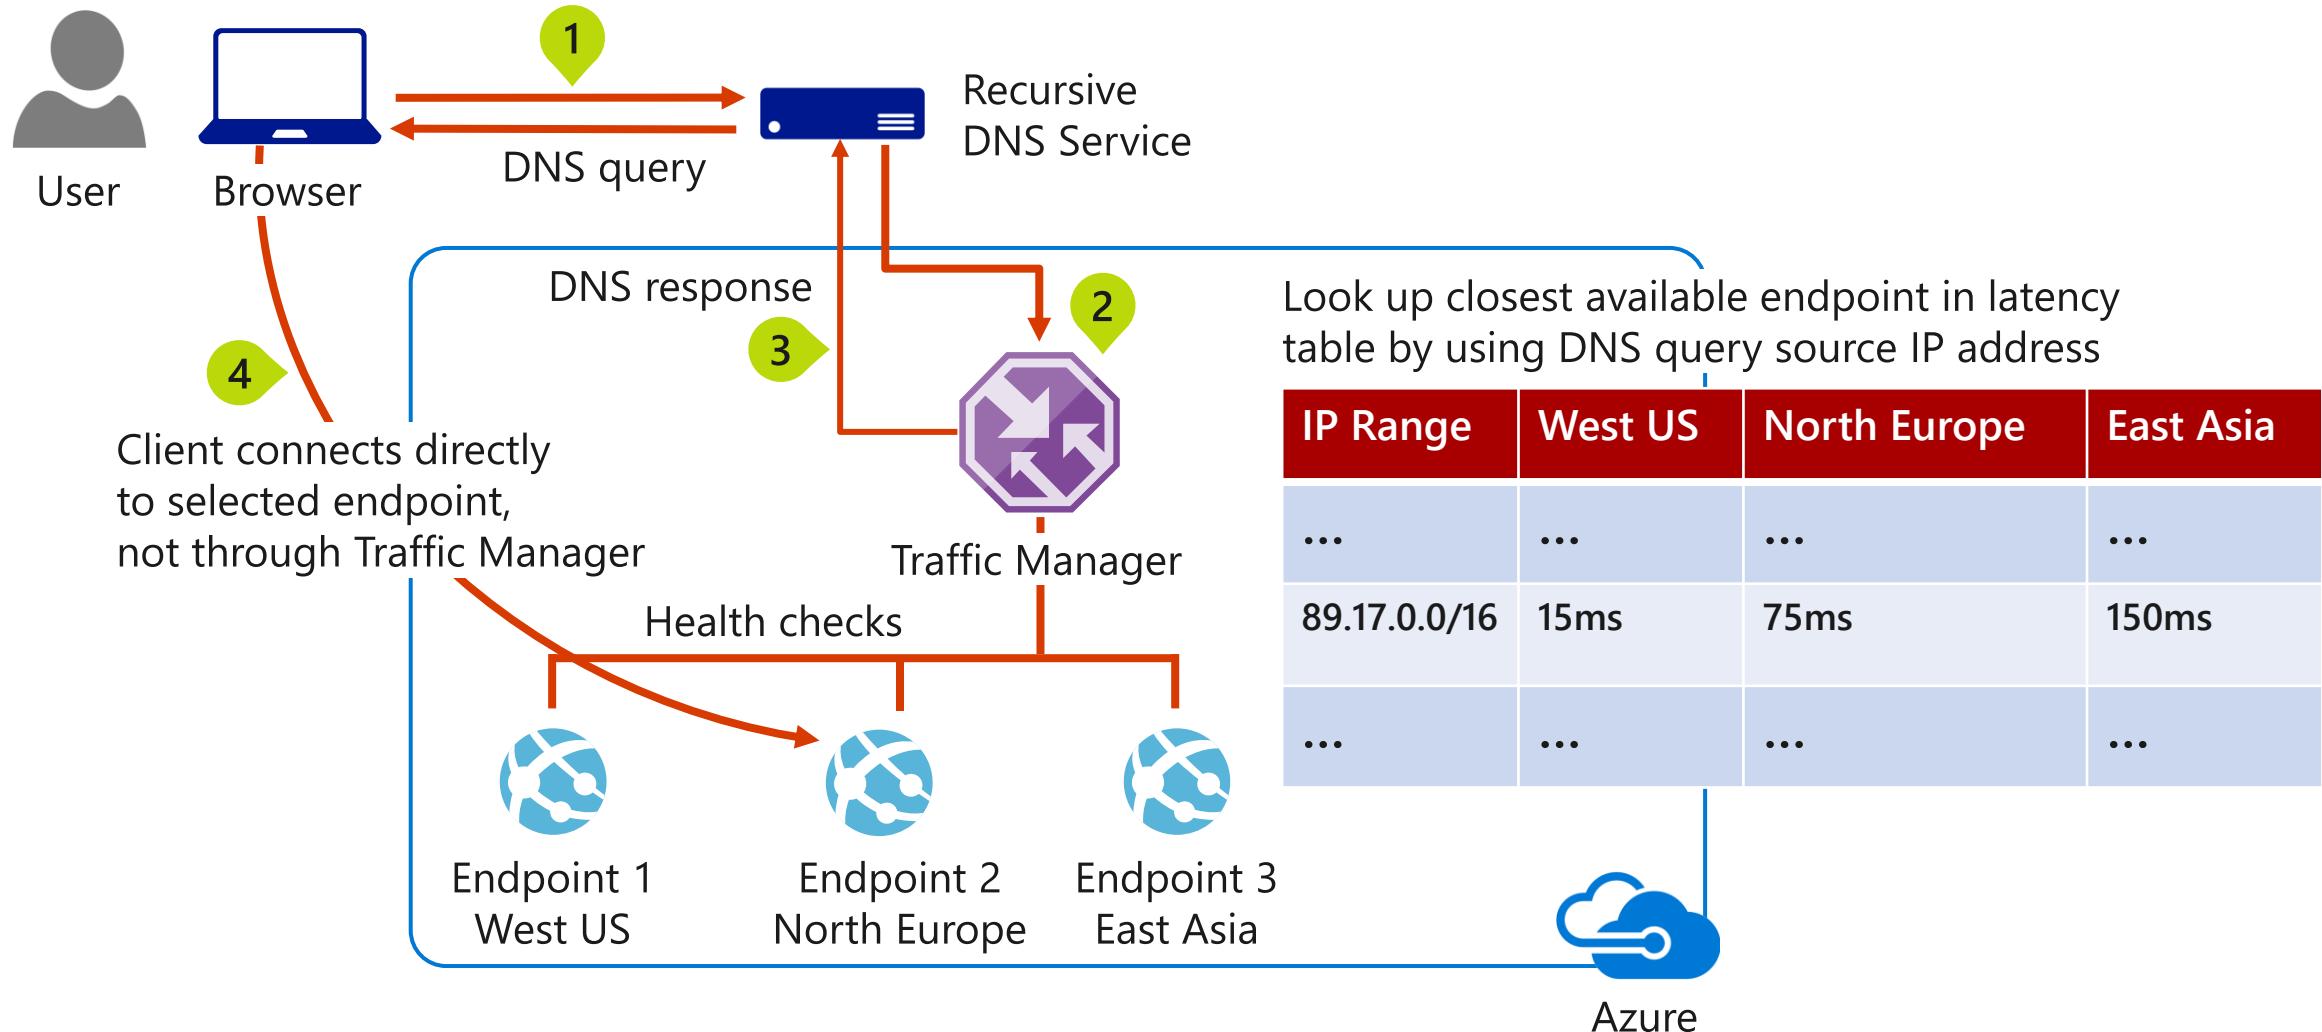
Priority traffic-routing method



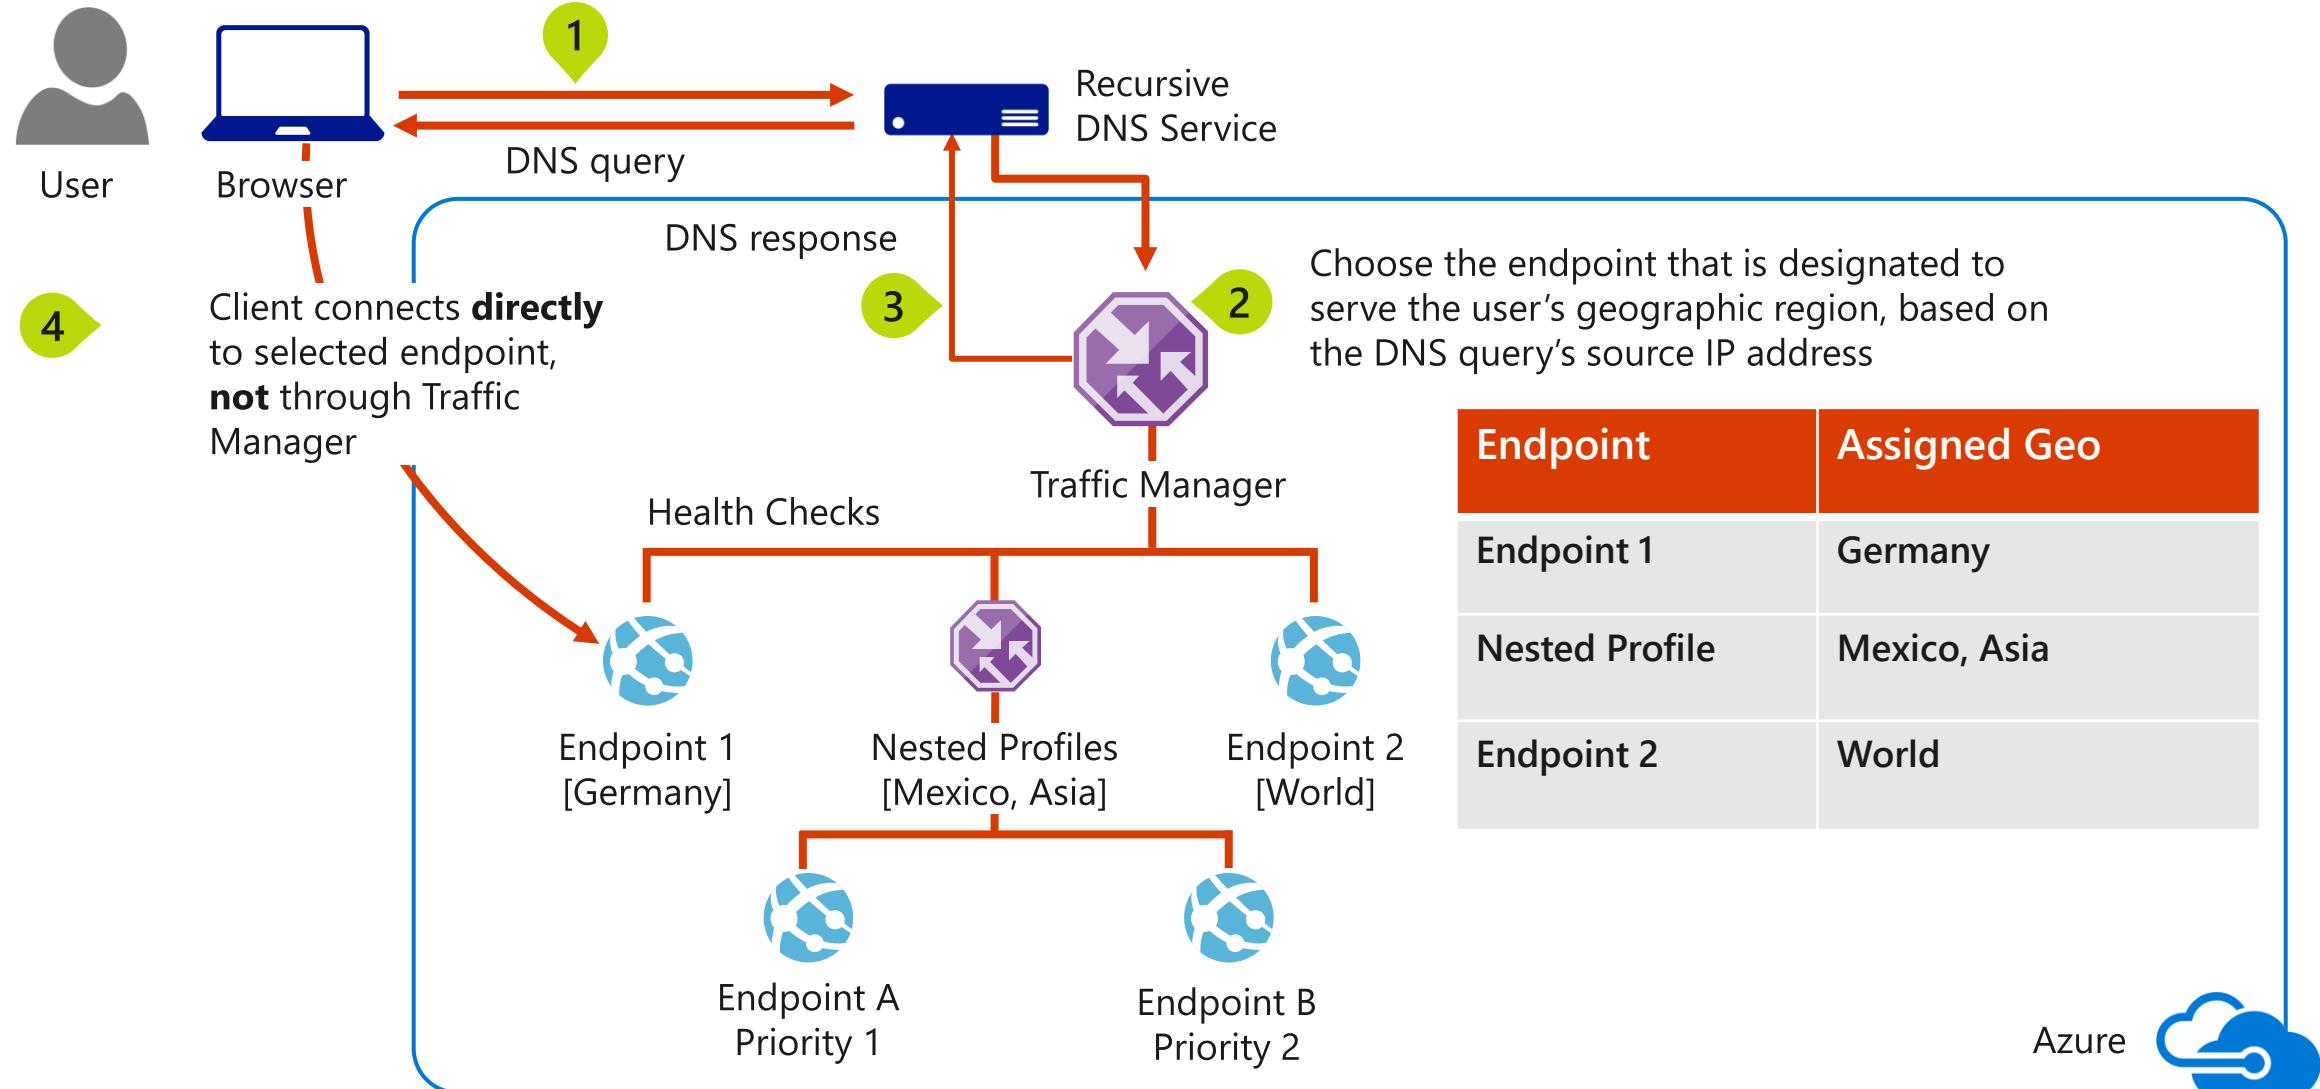
Weighted traffic-routing method



Performance traffic-routing method



Geographic traffic-routing method



Azure App Service Local Cache

- Provides a write-but-discard cache of your content
- Created asynchronously when site is started
- Automatically serves content once the cache is ready
- Faster than reading content from Azure Storage directly on each client request

App Service Environments (ASEs)

- App Service variant that provides a fully isolated and dedicated environment for securely running App Service apps at high scale
- Ideal for application workloads that require:
 - Very high scale, higher than typical App Service capacity
 - Network isolation and secure network access
 - High memory utilization
- Single or Multi-region
- Deployed to a virtual network
- An ASE is dedicated exclusively to a single subscription:
 - Max 100 instances

Lesson 04: Scaling App Service apps



Autoscale

- A primary advantage of the cloud is elastic scaling (the ability to use as much capacity as you need):
 - Scaling out as load increases
 - Scaling in when the extra capacity is not needed
- Many Microsoft Azure services provide the capability to scale both manually and automatically
- Autoscale refers to the capability of many of these services to monitor the application instances and automatically scale appropriately to handle the current usage of the application:
 - Using autoscale, your cloud service can scale out and in to exactly match the amount of instances needed for your specific computing pattern

Autoscale metrics

Metric	Metric identifier	Description
CPU	CpuPercentage	The average amount of CPU time used across all instances of the plan
Memory	MemoryPercentage	The average amount of memory used across all instances of the plan
Data in	BytesReceived	The average incoming bandwidth used across all instances of the plan
Data out	BytesSent	The average outgoing bandwidth used across all instances of the plan
HTTP queue	HttpQueueLength	The average number of both read and write requests that were queued on storage. A high disk queue length is an indication of an application that might be slowing down due to excessive disk I/O.
Disk queue	DiskQueueLength	The average number of HTTP requests that had to sit in the queue before being fulfilled. A high or increasing HTTP queue length is a symptom of a plan under a heavy load.

Autoscale patterns

- Scale based on CPU
- Scale differently on weekdays vs. weekends
- Scale differently during holidays
- Scale based on custom metric

Scale based on CPU

Scale differently on weekdays vs. weekends

Microsoft Azure Monitor - Autoscale > Autoscale setting

Autoscale setting WeekdayTrafficAsp (App Service plan)

Save Discard Disable autoscale

Configure Run history JSON Notify

Autoscale setting name: WeekdayTrafficAsp

Resource group: autoscaledemo

Instance count: 1

Default Auto created scale condition (edit) Delete

Scale mode: Scale based on a metric Scale to a specific instance count

Instance count: 3

Schedule: This scale condition is executed when none of the other scale condition(s) match

WeekendTraffic (edit) Delete

Scale mode: Scale based on a metric Scale to a specific instance count

Instance count: 1

Schedule: Specify start/end dates Repeat specific days

Repeat every: Monday Tuesday Wednesday Thursday Friday
 Saturday Sunday

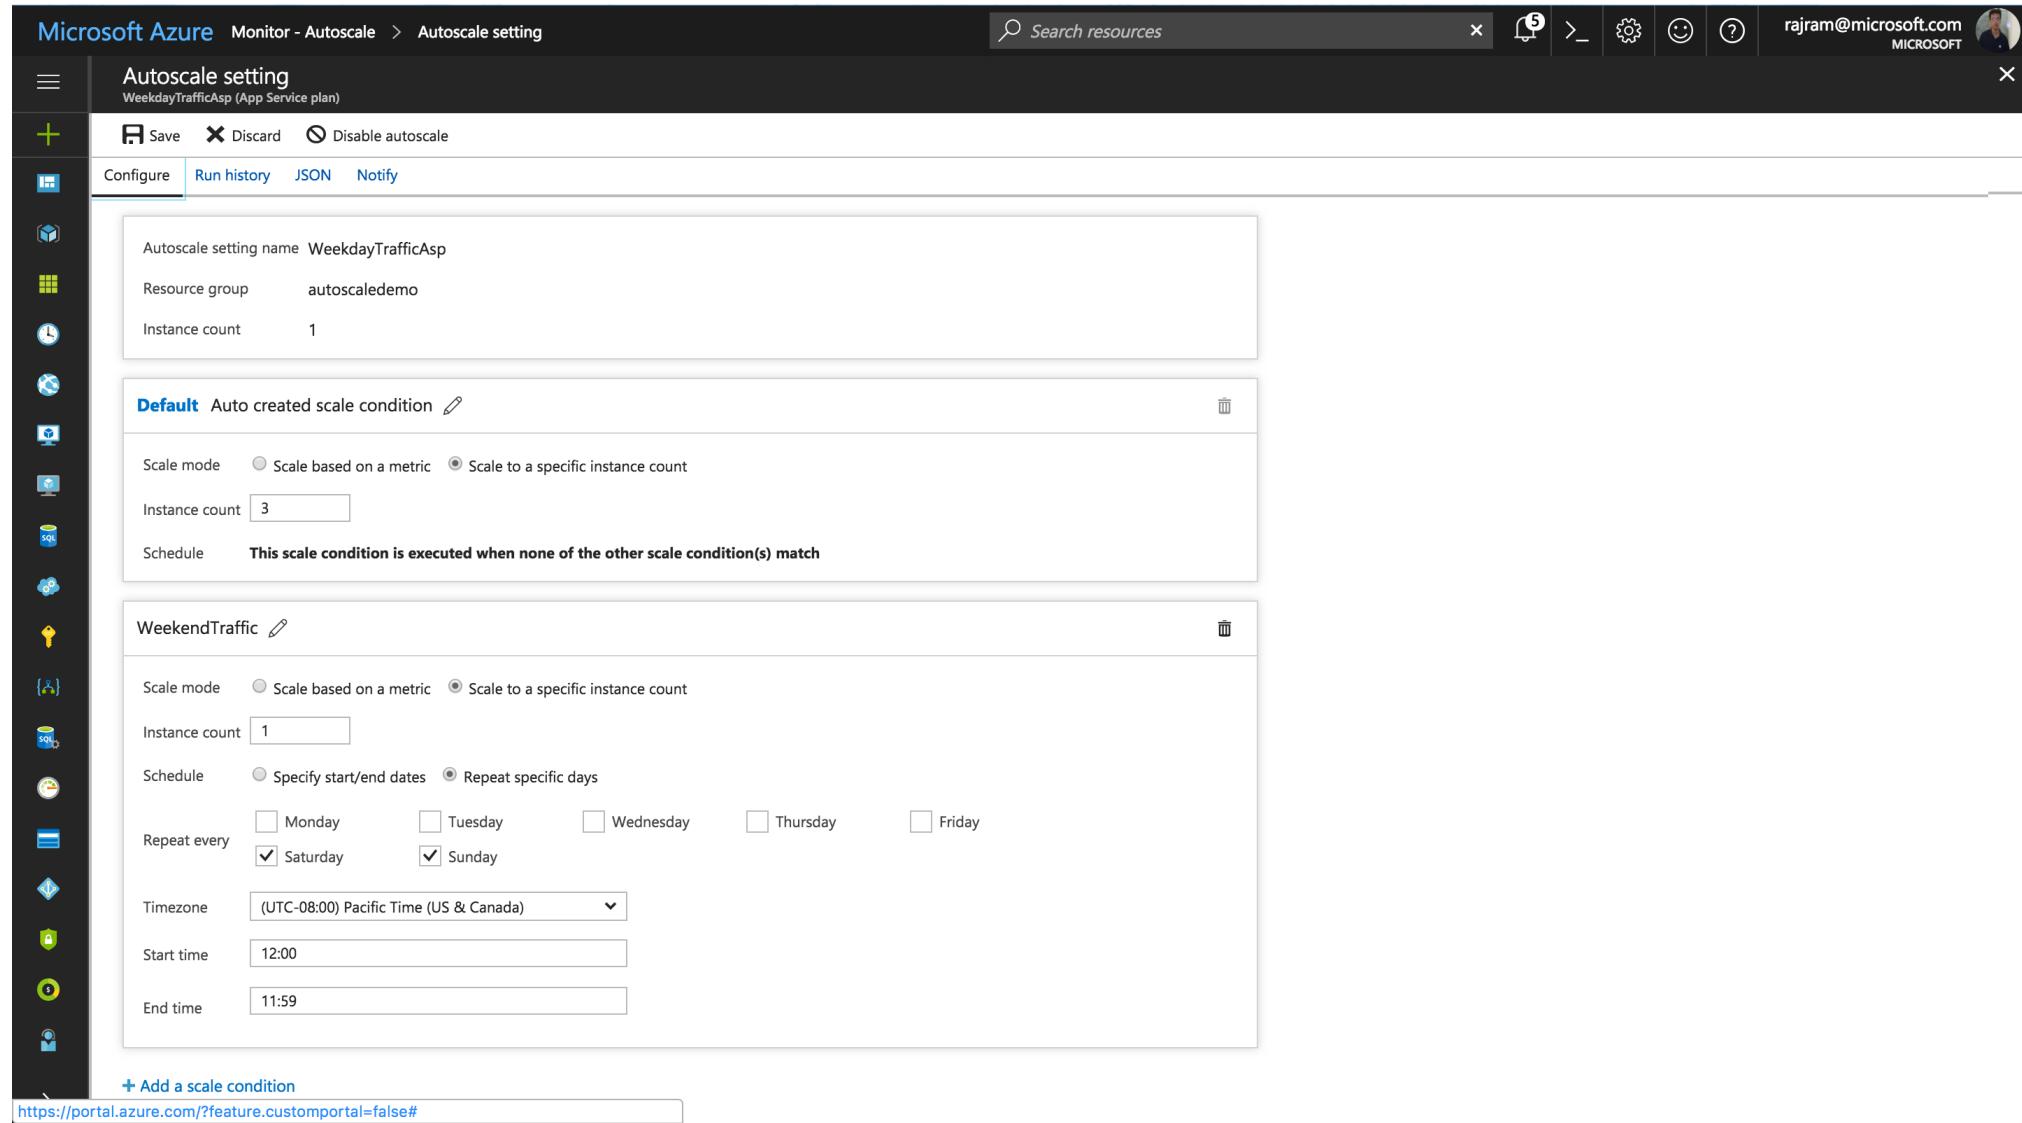
Timezone: (UTC-08:00) Pacific Time (US & Canada)

Start time: 12:00

End time: 11:59

[+ Add a scale condition](#)

<https://portal.azure.com/?feature.customportal=false#>



Scale differently during holidays

Microsoft Azure Monitor - Autoscale > Autoscale setting

Search resources X Notification bell Help Profile

Autoscale setting
HolidaySpikeAsp (App Service plan)

Save Discard Disable autoscale

Configure Run history JSON Notify

Autoscale setting name: HolidaySpikeAsp
Resource group: autoscaledemo
Instance count: 2

Default NormalScale edit

Scale mode: Scale based on a metric Scale to a specific instance count

Rules

Scale out

When: HolidaySpikeAsp (Average) CpuPercentage > 70 Increase instance count by 2

Scale in

When: HolidaySpikeAsp (Average) CpuPercentage < 30 Decrease instance count by 1

+ Add a rule

Instance limits: Minimum 2, Maximum 5, Default 2

Schedule: This scale condition is executed when none of the other scale condition(s) match

HolidayScale edit

Scale mode: Scale based on a metric Scale to a specific instance count

Instance count: 8

Schedule: Specify start/end dates Repeat specific days

Timezone: (UTC-08:00) Pacific Time (US & Canada)

Start date: 2017-11-30 23:00:00

End date: 2017-12-31 22:59:00

The screenshot shows the Microsoft Azure portal interface for managing autoscale settings. It displays two autoscale configurations: 'Default' (NormalScale) and 'HolidayScale'. The 'Default' configuration uses metrics to scale up (CpuPercentage > 70) and down (CpuPercentage < 30). The 'HolidayScale' configuration is set to a fixed instance count of 8 and applies from November 30 to December 31, 2017, in Pacific Time.

Scale based on custom metric

Microsoft Azure Monitor - Autoscale > Autoscale setting

Search resources (4) Configure Run history JSON Notify

Autoscale setting
contoso-web-api-asp (App Service plan)

Save Discard Disable autoscale

Configure Run history JSON Notify

Autoscale setting name Web api autoscale

Resource group contoso-web

Instance count 1

Default Auto created scale condition edit delete

Scale mode Scale based on a metric Scale to a specific instance count

Scale out

When	loans-app-ai	(Total) customMetrics/LoanSubmissions > 100	Increase instance count by 1
------	--------------	---	------------------------------

Rules

Scale in

When	loans-app-ai	(Total) customMetrics/LoanSubmissions < 25	Decrease instance count by 1
------	--------------	--	------------------------------

+ Add a rule

Instance limits

Minimum 1	Maximum 5	Default 2
------------------------	------------------------	------------------------

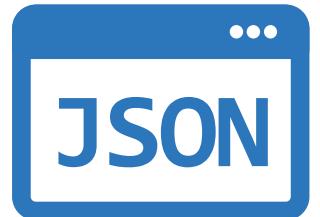
Schedule **This scale condition is executed when none of the other scale condition(s) match**

+ Add a scale condition

The screenshot shows the Azure portal interface for managing autoscale settings. On the left, there's a sidebar with various service icons. The main area is titled 'Autoscale setting' for a resource named 'contoso-web-api-asp'. It displays basic details like the name, resource group, and current instance count (1). Below this, a 'Default' scale condition is shown, which uses a custom metric ('(Total) customMetrics/LoanSubmissions') to trigger scaling. There are two rules: one to scale out when the metric value is greater than 100 (increasing instances by 1), and another to scale in when it's less than 25 (decreasing instances by 1). The instance limits are set from 2 to 5. A note at the bottom states that this condition runs when no other conditions match.

Autoscale setting schema

```
{  
  "id": "...", "name": "demoSetting", "type": "Microsoft.Insights/autoscaleSettings",  
  "location": "East US",  
  "properties": {  
    "enabled": true, "targetResourceUri": "...",  
    "profiles": [  
      {  
        "name": "mainProfile",  
        "capacity": {  
          "minimum": "1", "maximum": "4", "default": "1"  
        },  
        "rules": [ ... ]  
      }  
    ]  
  }  
}
```



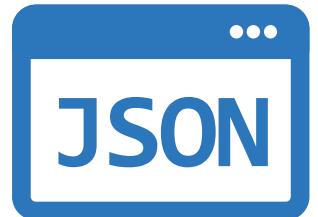
Autoscale setting schema – scale-out rule

```
{  
    "metricTrigger": {  
        "metricName": "Percentage CPU", "metricResourceUri": "...",  
        "timeGrain": "PT1M",  
        "statistic": "Average",  
        "timeWindow": "PT10M",  
        "timeAggregation": "Average",  
        "operator": "GreaterThan",  
        "threshold": 85  
    "scaleAction": {  
        "direction": "Increase", "type": "ChangeCount",  
        "value": "1",  
        "cooldown": "PT5M"  
}
```



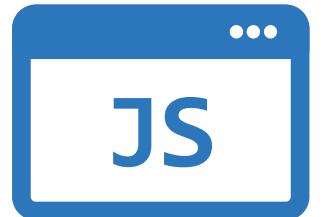
Autoscale setting schema – scale-in rule

```
{  
    "metricTrigger": {  
        "metricName": "Percentage CPU", "metricResourceUri": "...",  
        "timeGrain": "PT1M",  
        "statistic": "Average",  
        "timeWindow": "PT10M",  
        "timeAggregation": "Average",  
        "operator": "LessThan",  
        "threshold": 60  
    "scaleAction": {  
        "direction": "Decrease", "type": "ChangeCount",  
        "value": "1",  
        "cooldown": "PT5M"  
}
```



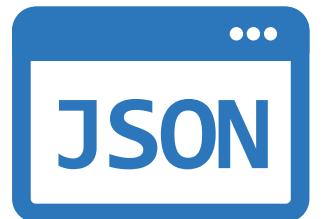
Autoscale setting schema – fixed date profile

```
"fixedDate": {  
    "timeZone": "Pacific Standard Time",  
    "start": "2020-12-26T00:00:00",  
    "end": "2020-12-26T23:59:00"  
}
```



Autoscale setting schema – weekdays and weekends

```
"recurrence": {  
    "frequency": "Week",  
    "schedule": {  
        "timeZone": "Pacific Standard Time",  
        "days": [  
            "Saturday"  
        ],  
        "hours": [  
            0  
        ],  
        "minutes": [  
            0  
        ]  
    }  
}
```

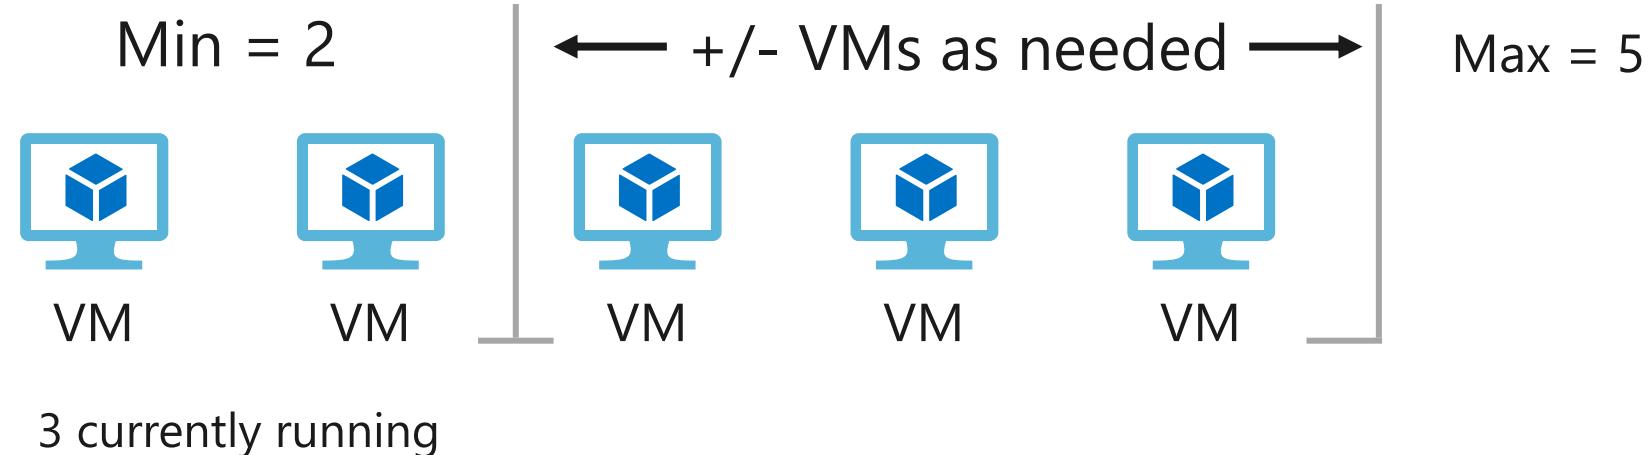


Autoscale concepts

- Each resource can have one autoscale setting:
 - Autoscale settings can have one-to-many profiles
 - Profiles can have one-to-many rules
- Autoscale increases instances horizontally within bounds:
 - Bounds are set by using the minimum, maximum, and default values
- Thresholds are calculated at an instance level
- Autoscale successful actions and failures are logged to the Activity Log

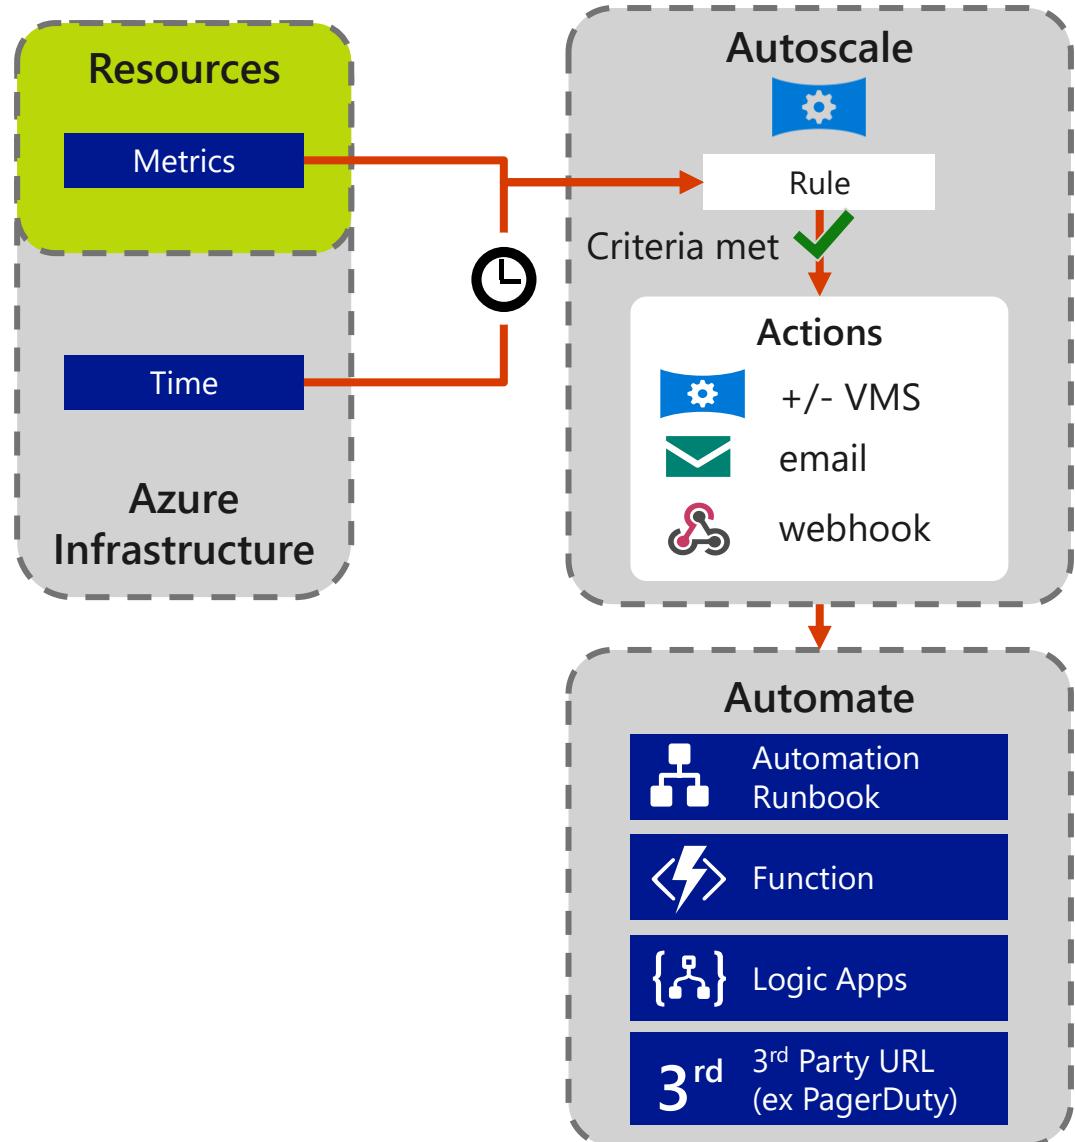
Autoscale thresholds

- Scale is constrained to a minimum and maximum:
 - Your current instance count must be between the minimum and maximum:
 - Minimum can help guarantee availability
 - Maximum can help control costs



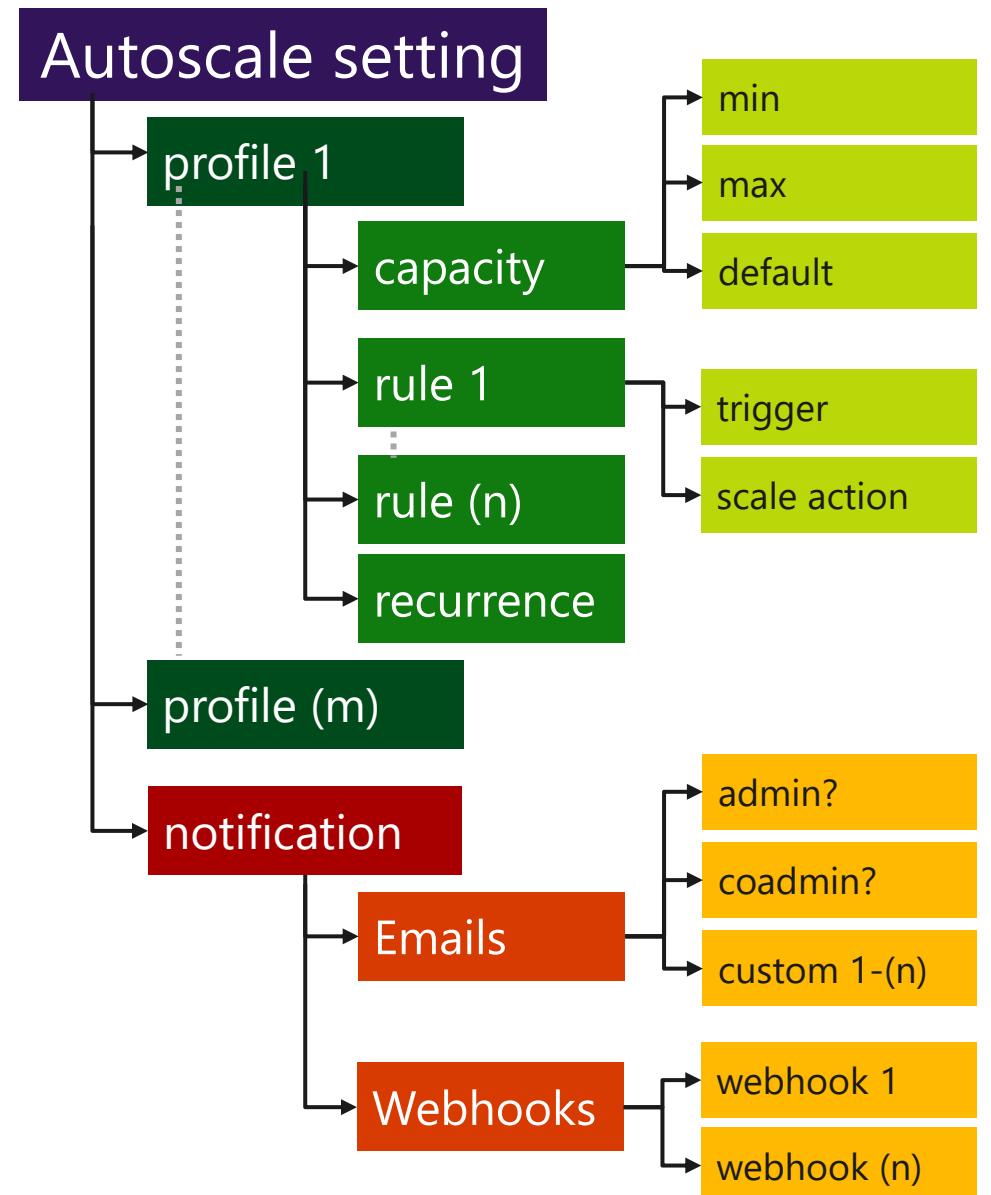
Autoscale workflow

1. Metrics are measured for a resource
2. When conditions are met (threshold surpassed), autoscale triggers:
 - a. Perform scale actions
 - b. Send notifications (alerts)
 - c. Send messages to webhooks for external automation



Autoscale hierarchy

- One autoscale setting
- Settings have one or more profiles
- Profiles have one or more rules:
 - Profiles can also have recurrences and capacity settings
- Notifications can be directly associated with an autoscale setting



Best practices

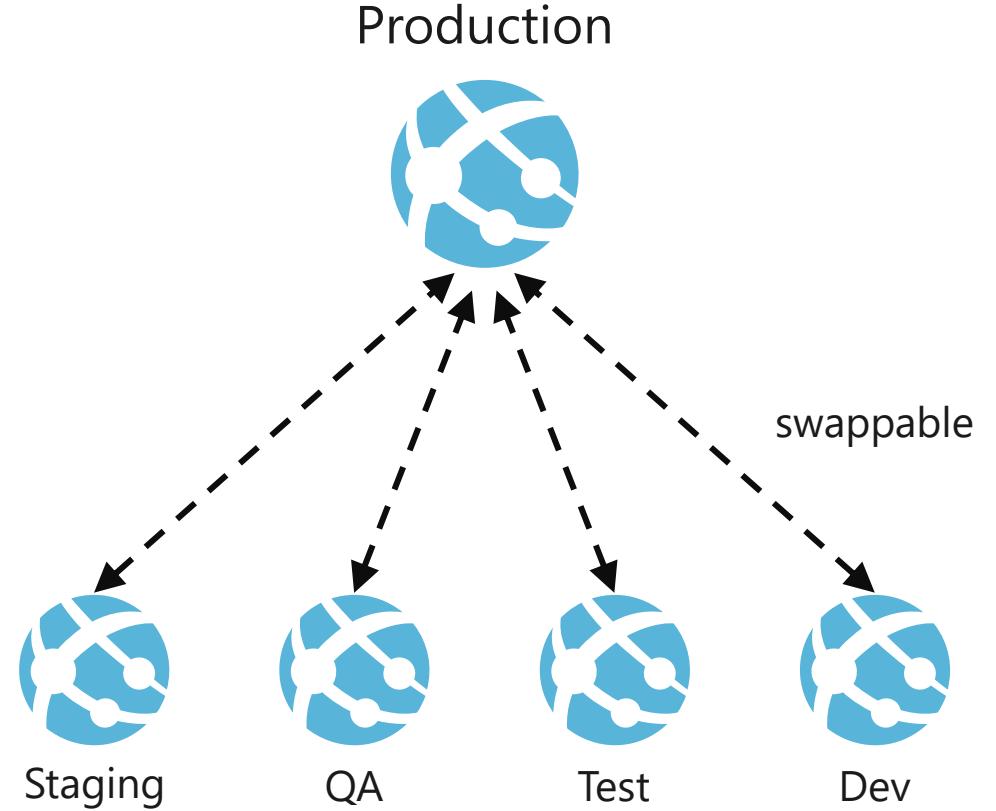
- Ensure that the maximum and minimum values are different and have an adequate margin between them
- Manual scaling is reset by autoscale min and max
- Always use a scale-out and scale-in rule combination that performs an increase and decrease
- Choose the appropriate statistic for your diagnostics metric
- Choose the thresholds carefully for all metric types

Lesson 05: Azure App Service staging environments

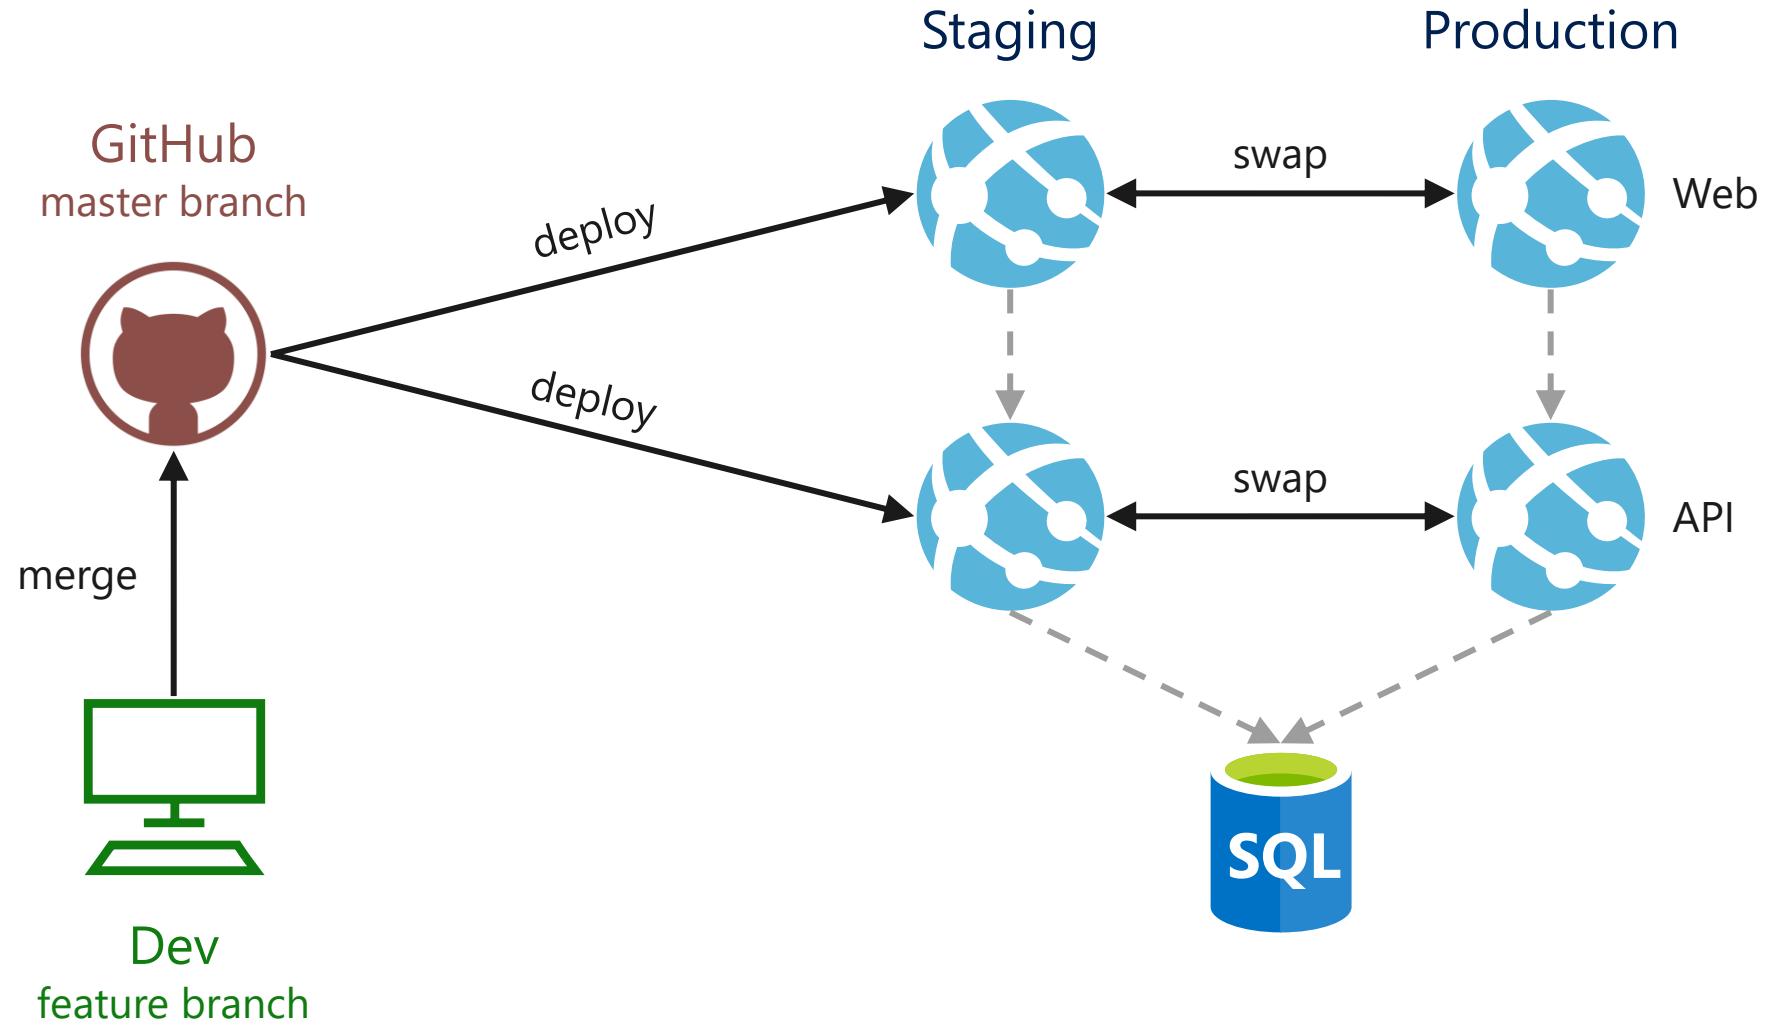


Deployment slots

- Live apps with their own:
 - Host names
 - Content
 - Configuration
- Can be swapped between each other.
 - Staging ↔ Production
 - Production ↔ Staging
 - Dev ↔ Test
 - Test ↔ QA
 - QA ↔ Staging



Modern deployment workflow

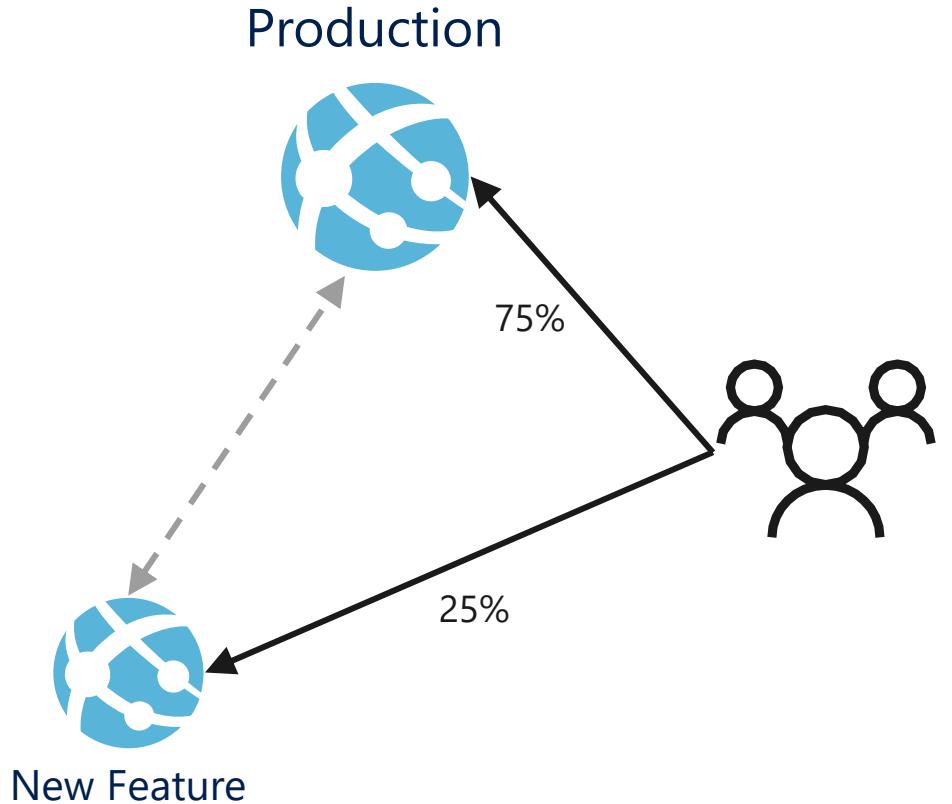


Auto swap

- Automatically swaps an application to the production slot:
 - Performed after the application is "warmed up"
 - Swaps the deployment target slot with the production slot
- Deploy apps continuously while minimizing cold starts and downtime

Route traffic between slots

- All traffic is normally routed to production:
 - Production slot has 100% weighting
- You can manually configure the weight of traffic between multiple slots



Automate slot management - Azure PowerShell

```
New-AzWebAppSlot -ResourceGroupName "<groupname>" -Name "<appname>" -Slot "staging" -  
AppServicePlan "<appserviceplanname>"
```

Create new slot

Target

```
$Params = @{targetSlot = "[slot name - e.g. "production"]"}
```

```
Invoke-AzResourceAction -ResourceGroupName "<groupname>"  
-ResourceType "Microsoft.Web/sites/slots" -ResourceName "<appname>/staging"  
-Action "slotsswap" -Parameters $Params -ApiVersion 2015-07-01
```

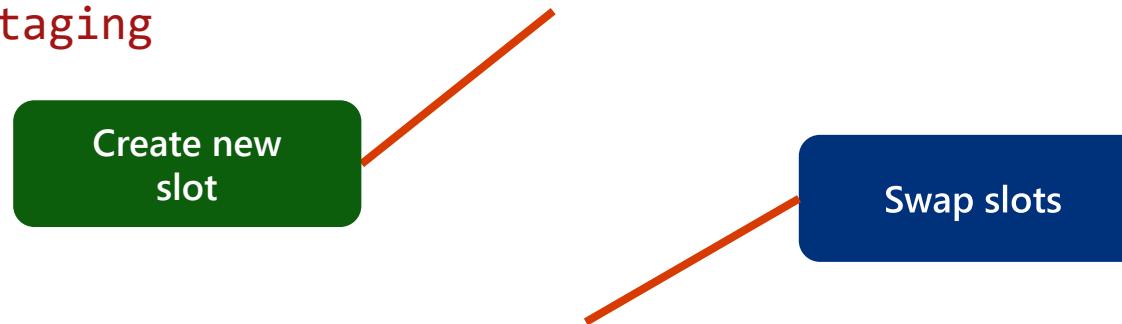
Swap slots

Source

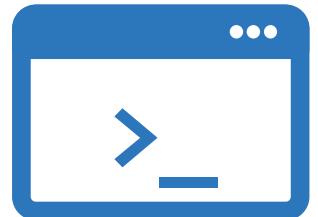
PS

Automate slot management - Azure CLI

```
az webapp deployment slot create --name <appname> --resource-group <groupname> --slot staging
```



```
az webapp deployment slot swap --resource-group <groupname> --name <appname> --slot staging --target-slot production
```



Lab: Building a web application on Azure platform as a service offerings

Duration



Lab sign-in information

